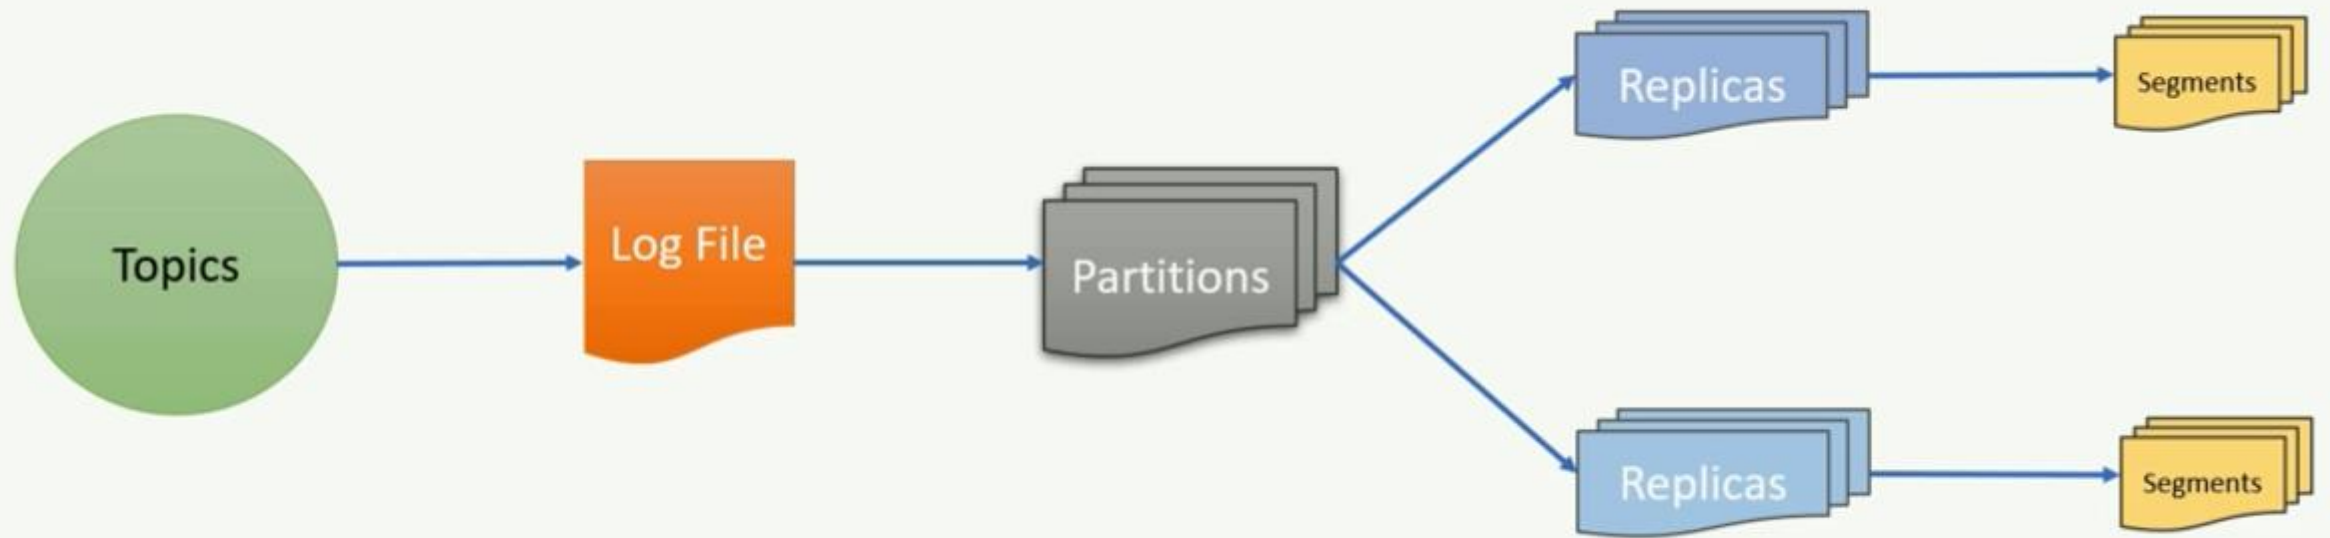




# **Kafka Storage System**





Topic  
(Logical Name)

Invoice

Invoice  
P0

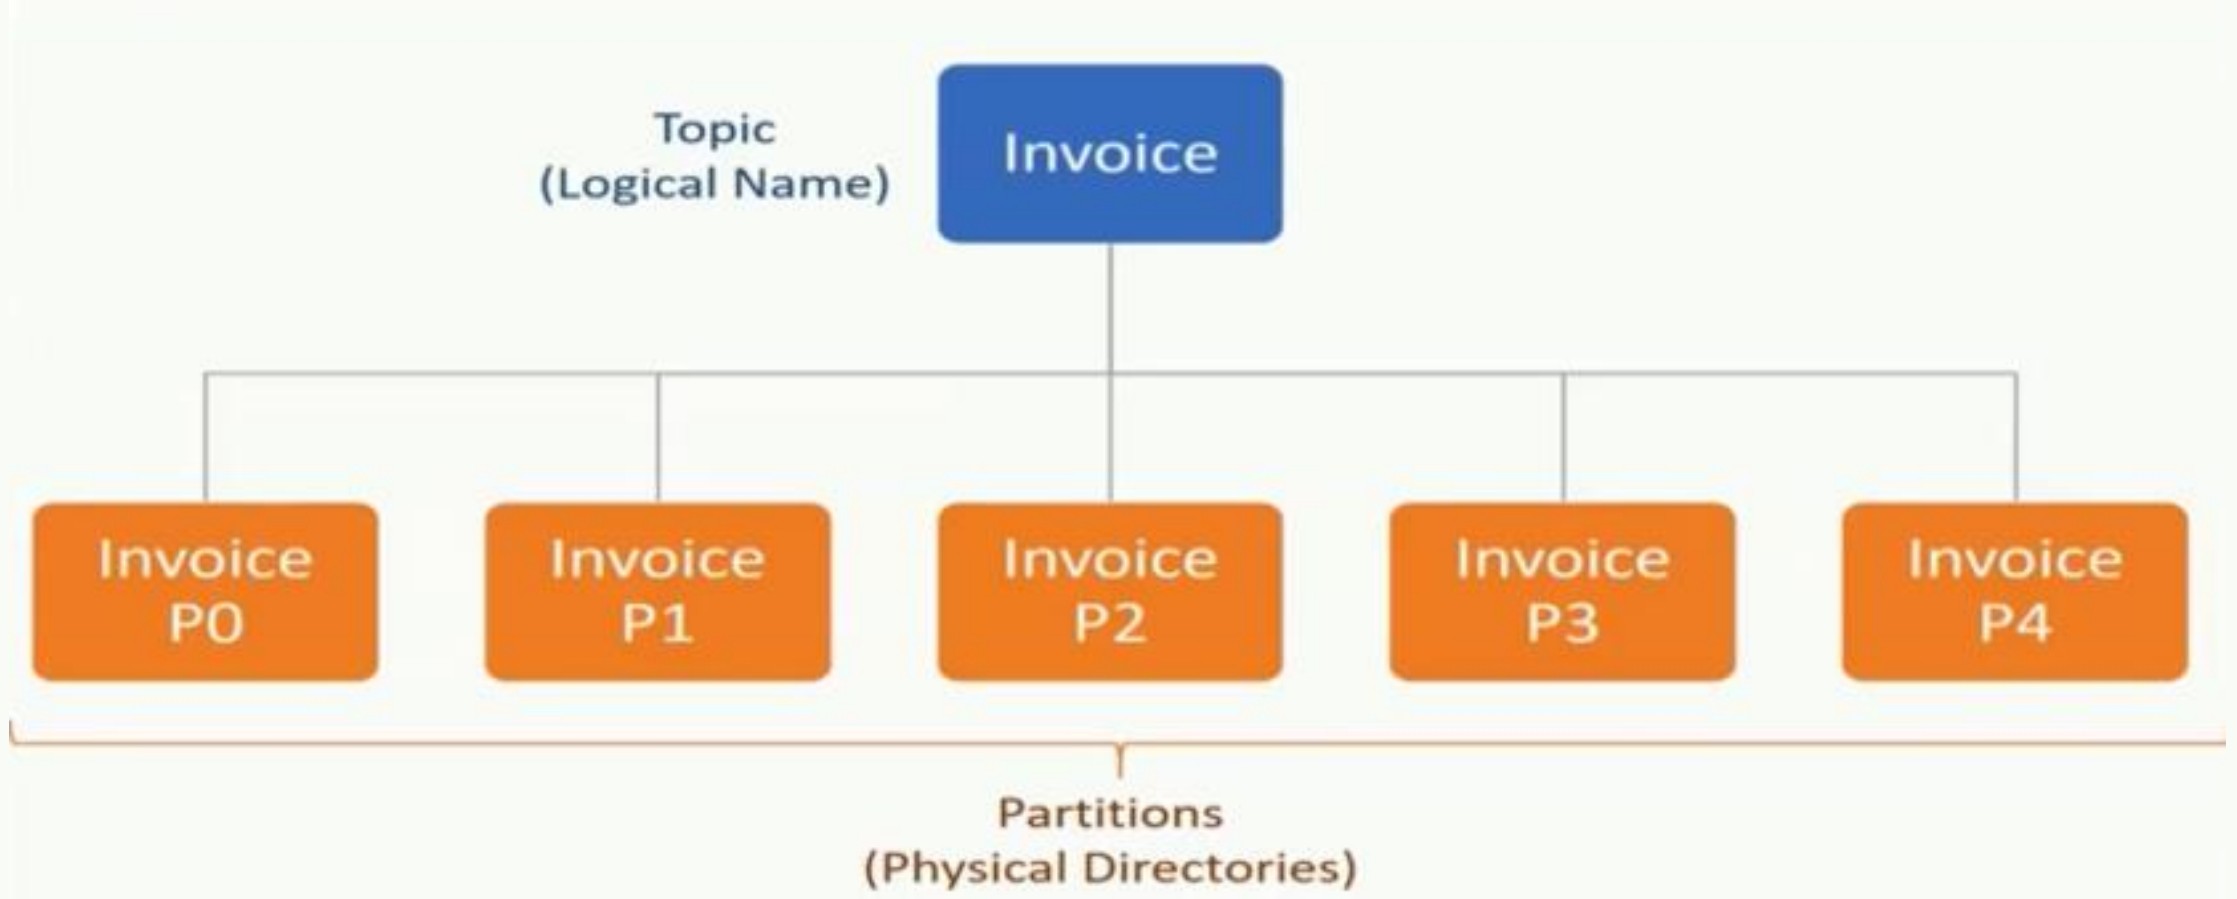
Invoice  
P1

Invoice  
P2

Invoice  
P3

Invoice  
P4

Partitions  
(Physical Directories)

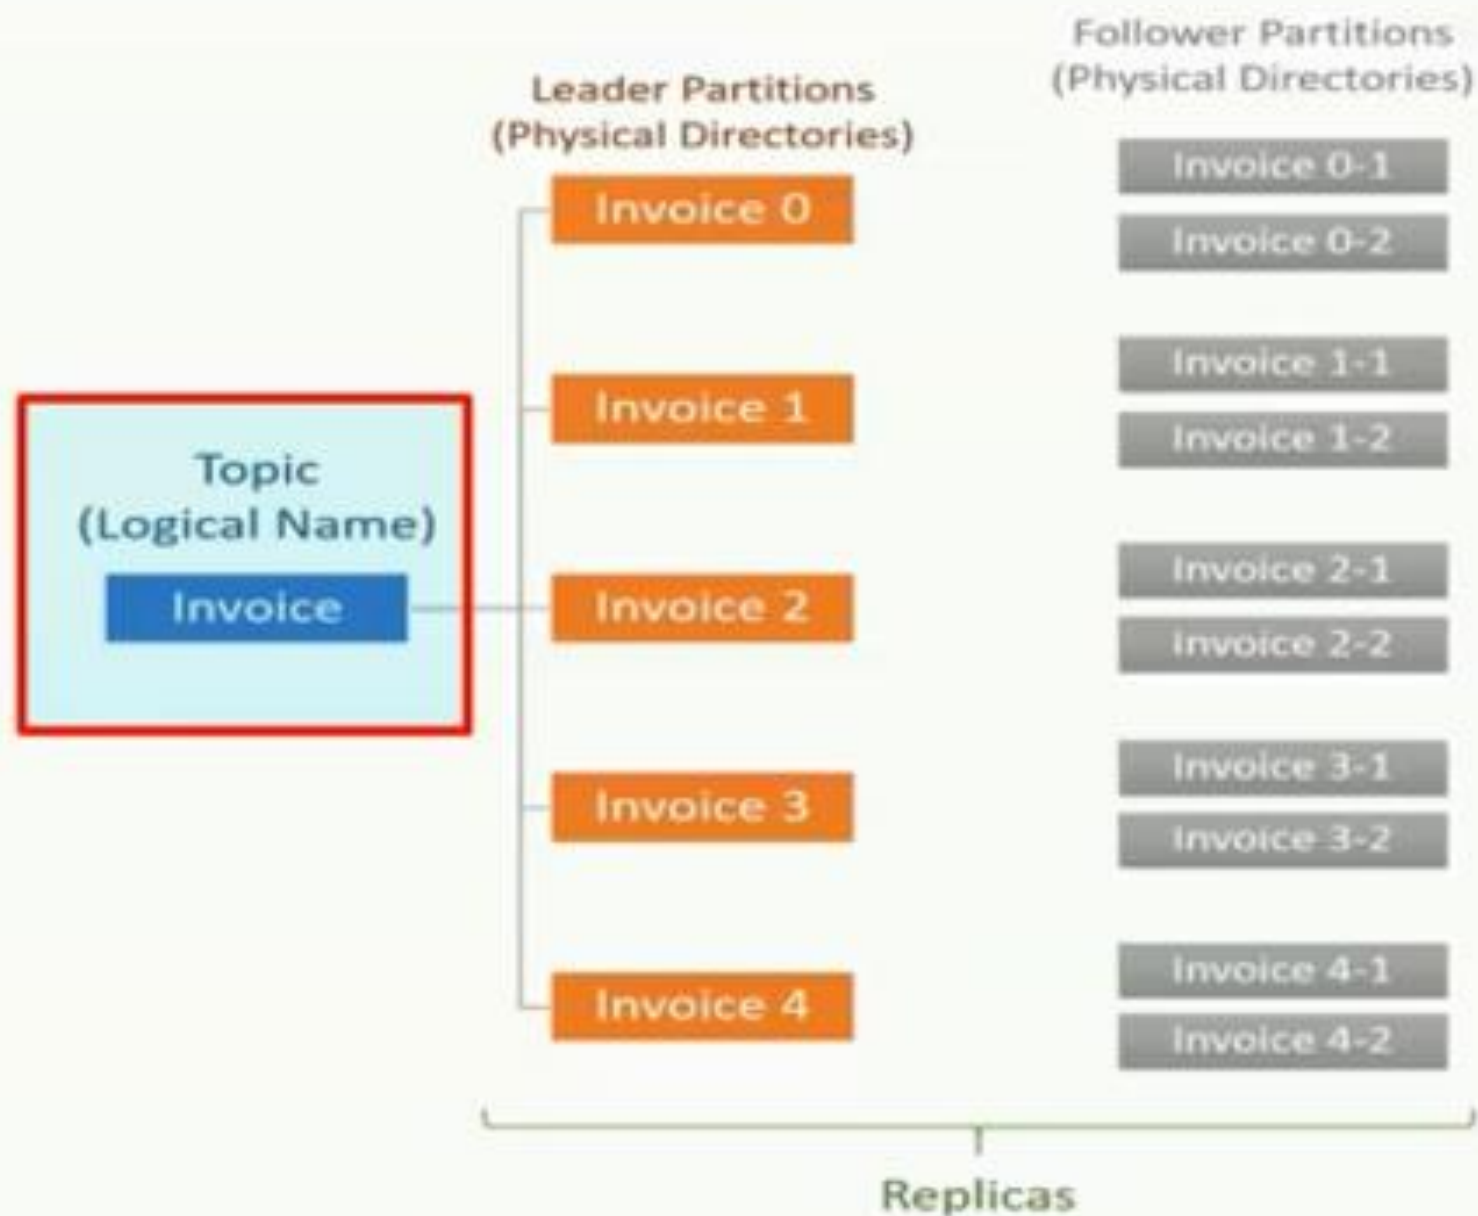


# What is Replication Factor?


Number of Copies for each Partition



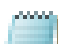




Number of Replicas (15) = Partitions (5) X Replication (3)

1. Leader Partitions
2. Follower Partitions















```
amin@DESKTOP-3F1FI9I:~$ kafkacat -b localhost:9092 -L -t numbers
Metadata for numbers (from broker 1: localhost:9092/1):
 3 brokers:
   broker 1 at localhost:9092 (controller)
   broker 2 at localhost:9094
   broker 3 at localhost:9096
 1 topics:
   topic "numbers" with 4 partitions:
     partition 0, leader 3, replicas: 3, isrs: 3
     partition 1, leader 1, replicas: 1, isrs: 1
     partition 2, leader 2, replicas: 2, isrs: 2
     partition 3, leader 1, replicas: 1, isrs: 1
```



 00000000000000000000.index	5/27/2024 12:17 PM	INDEX File	1 KB
 00000000000000000000.log	5/27/2024 12:15 PM	Text Document	19 KB
 00000000000000000000.timeindex	5/27/2024 12:17 PM	TIMEINDEX File	1 KB
 000000000000000000006.snapshot	5/27/2024 12:00 PM	SNAPSHOT File	1 KB
 0000000000000000000085.snapshot	5/27/2024 12:17 PM	SNAPSHOT File	1 KB
 leader-epoch-checkpoint	5/27/2024 12:05 PM	File	1 KB
 partition.metadata	5/27/2024 11:46 AM	METADATA File	1 KB



 0000000000000000000217.index  
 0000000000000000000217  
 0000000000000000000217.timeindex  
 00000000000000000004411.index  
 00000000000000000004411  
 00000000000000000004411.timeindex  
 00000000000000000008534.index  
 00000000000000000008534  
 00000000000000000008534.timeindex  
 00000000000000000012664.index  
 00000000000000000012664  
 00000000000000000012664.timeindex

**1 GB = 1073741824 bytes**

```
# The minimum age of a log file to be eligible for deletion due to age
```

```
log.retention.hours=168
```

```
# A size-based retention policy for logs. Segments are pruned from the log unless the remaining  
# segments drop below log.retention.bytes. Functions independently of log.retention.hours.
```


```
#log.retention.bytes=1073741824
```

```
# The maximum size of a log segment file. When this size is reached a new log segment will be cr
```

```
#log.segment.bytes=1073741824
```


```
# The interval at which log segments are checked to see if they can be deleted according  
# to the retention policies
```

```
log.retention.check.interval.ms=300000
```





Every offset stores in a 64-bit integer numbers that its range is

000000000000000000000000 till  
18/446/744/073/709/551/615



1. Topic Name
2. Partition Number
3. Offset Number

```
SELECT * FROM invoice WHERE customer-id = 091
```



- 00000000000000000000217.index
- 00000000000000000000217
- 00000000000000000000217.timeindex
- 000000000000000000004411.index
- 000000000000000000004411
- 000000000000000000004411.timeindex
- 000000000000000000008534.index
- 000000000000000000008534
- 000000000000000000008534.timeindex
- 0000000000000000000012664.index
- 0000000000000000000012664
- 0000000000000000000012664.timeindex

Upon the request of clients Kafka with use of **.index** file can locate the offset which has been requested.

Clients can request messages based on time. So Kafka uses **.timeindex**.



Kafka stores records in **bytes**, and when a fetch request comes in from a consumer, Kafka returns records in bytes.

The screenshot shows the Apache Kafka Admin Console interface. On the left, the 'Clusters' tree is expanded to show '16-kafka-03' > 'Topics' > 'invoice' > 'Partitions' > 'Partition 0'. Below this, the 'Consumers' section is highlighted with a red box, showing 'console-consumer-41661'. The main panel displays the 'Data' tab for 'Partition 0', showing a table of records. The table has three columns: 'Offset', 'Key', and 'Value'. The 'Value' column contains hexadecimal strings. A red box highlights the 'Value' column in the table.

Offset	Key	Value
0	00000001	53696D706C65204D6573736167652D31
1	0000001B	53696D706C65204D6573736167652D3237
2	0000001E	53696D706C65204D6573736167652D3330
3	00000020	53696D706C65204D6573736167652D3332
4	00000021	53696D706C65204D6573736167652D3333
5	00000025	53696D706C65204D6573736167652D3337
6	00000031	53696D706C65204D6573736167652D3439
7	00000032	53696D706C65204D6573736167652D3530
8	00000039	53696D706C65204D6573736167652D3537
9	0000003E	53696D706C65204D6573736167652D3632
10	00000046	53696D706C65204D6573736167652D3730
11	0000004B	53696D706C65204D6573736167652D3735
12	00000051	53696D706C65204D6573736167652D3831
13	00000059	53696D706C65204D6573736167652D3839
14	0000005B	53696D706C65204D6573736167652D3931
15	0000005D	53696D706C65204D6573736167652D3933
16	0000005F	53696D706C65204D6573736167652D3935
17	00000064	53696D706C65204D6573736167652D4000