

Introduction to windowing in the world of Streaming

By Apache Spark



1. Stateless transformations
2. Stateful transformations

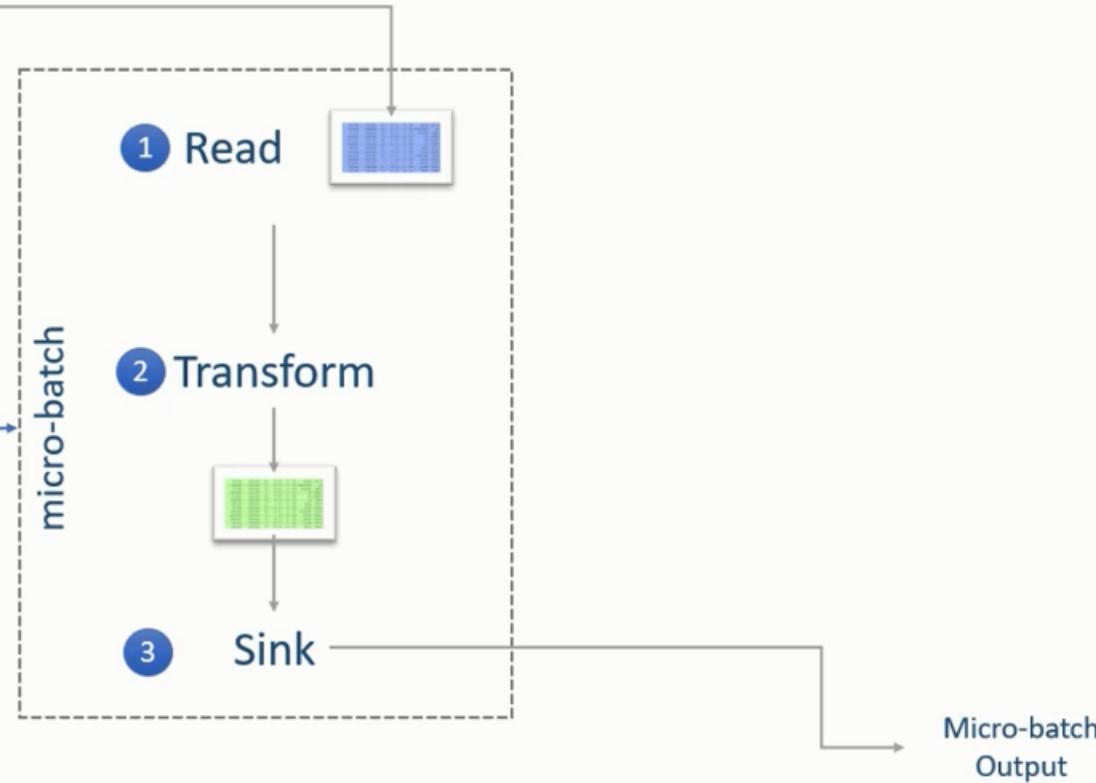


What is the state?



Streaming
Source

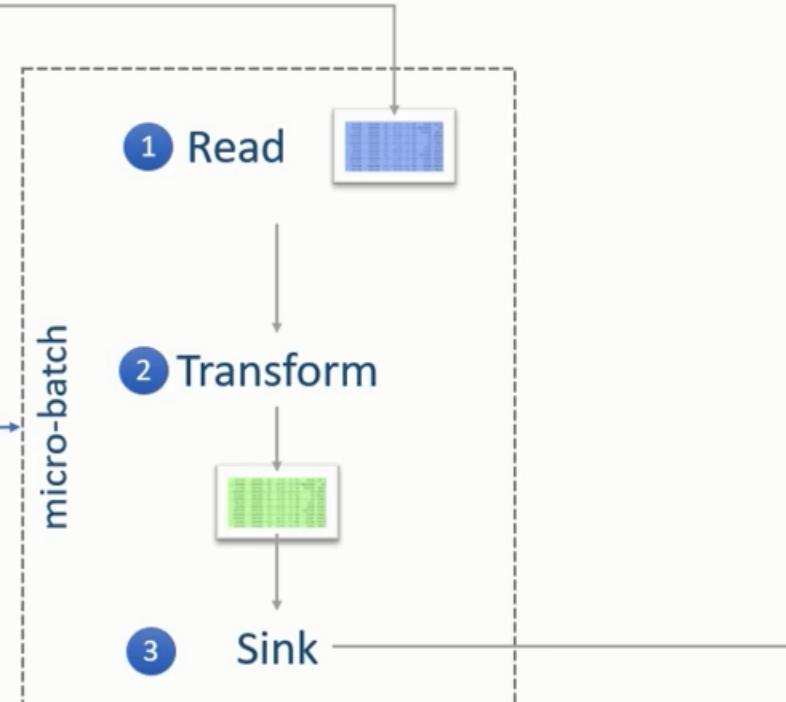
BG Thread





Streaming
Source

BG Thread



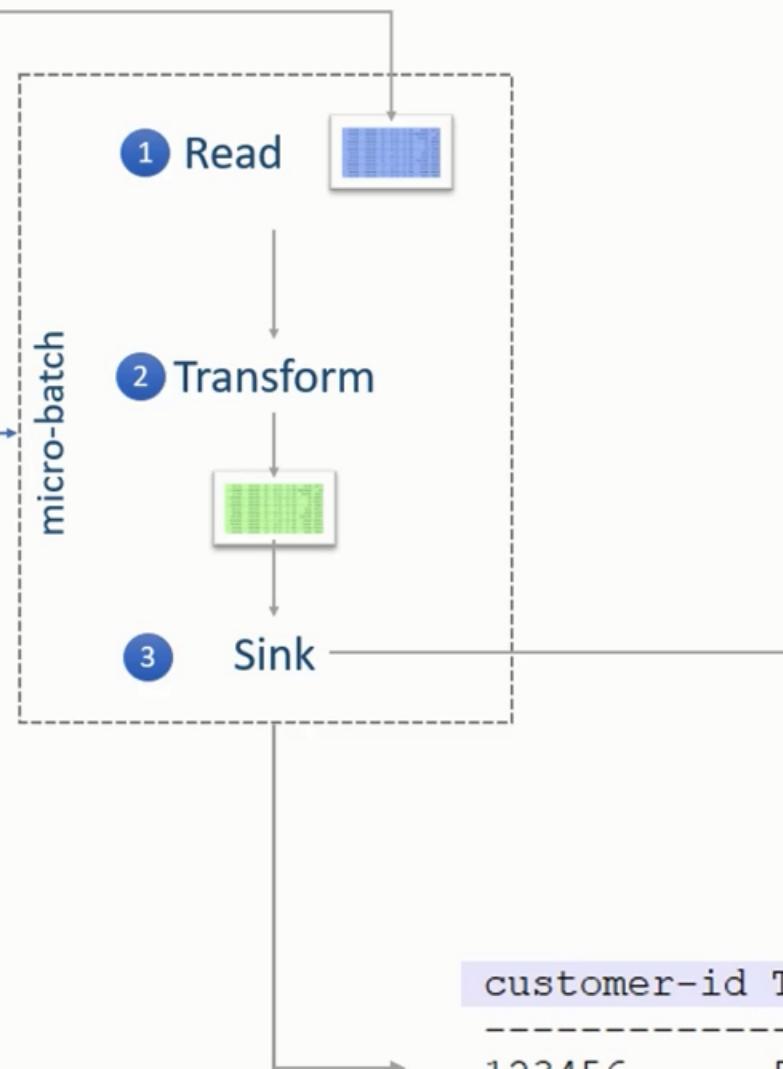
customer-id	Invoice Amount	Total Purchase	TotalRewards
123456	500	500	100

Micro-batch
Output



Streaming
Source

BG Thread



First purchase

customer-id	Invoice Amount	Total Purchase	TotalRewards
123456	500	500	100

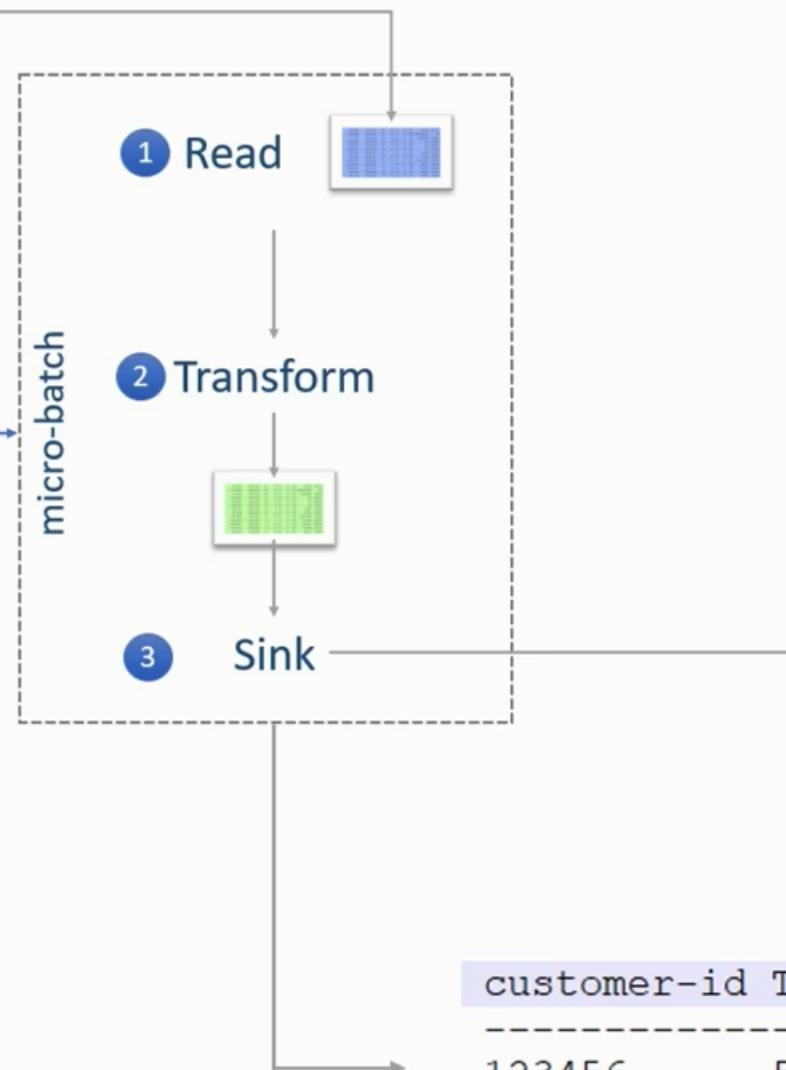
customer-id	Total Purchase	TotalRewards
123456	500	100

State Store



Streaming
Source

BG Thread



Second purchase

customer-id	Invoice Amount	Total Purchase	TotalRewards
123456	300	800	160

Micro-batch
Output

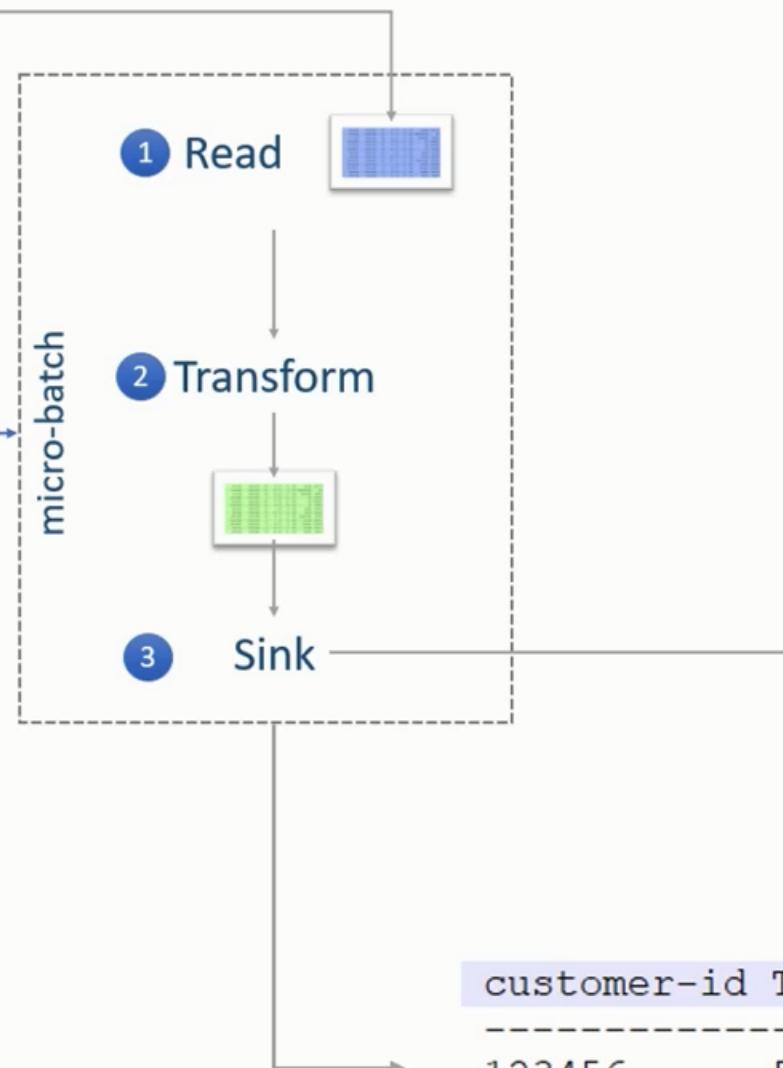
customer-id	Total Purchase	TotalRewards
123456	500	100

State Store



Streaming
Source

BG Thread



Second purchase

customer-id	Invoice Amount	Total Purchase	TotalRewards
123456	300	800	160

Micro-batch
Output

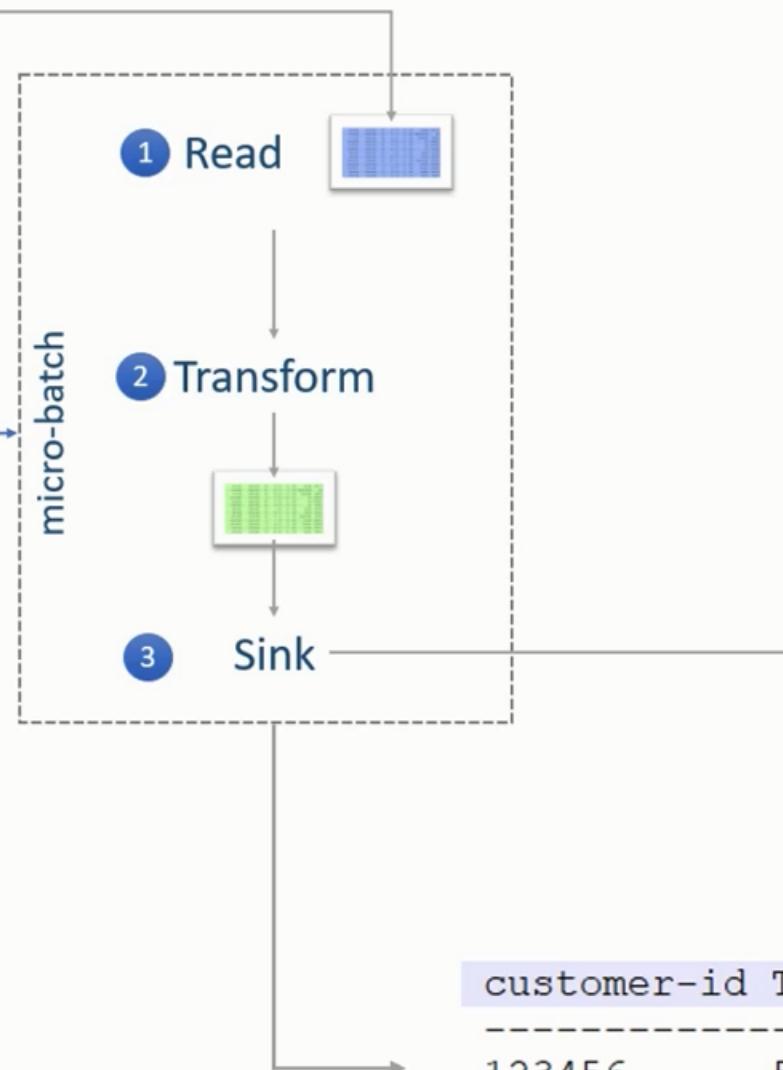
customer-id	Total Purchase	TotalRewards
123456	500	100

State Store



Streaming
Source

BG Thread



Second purchase

customer-id	Invoice Amount	Total Purchase	TotalRewards
123456	300	800	160

Micro-batch
Output

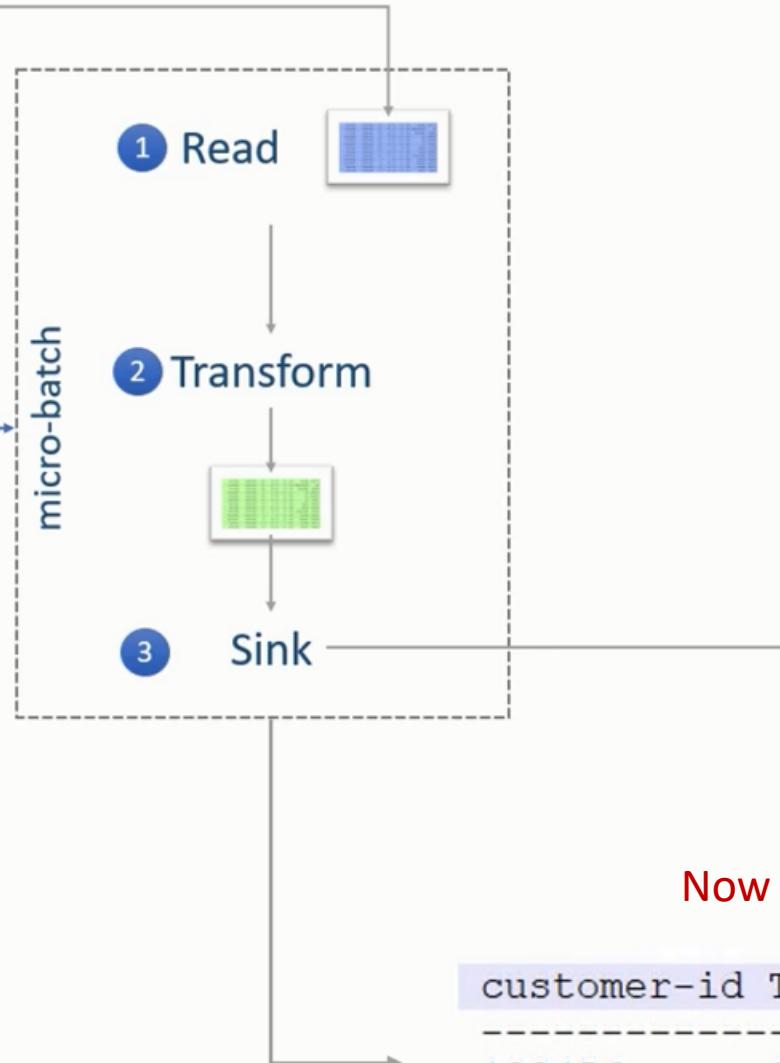
customer-id	Total Purchase	TotalRewards
123456	500	100

State Store



Streaming
Source

BG Thread



Second purchase

customer-id	Invoice Amount	Total Purchase	TotalRewards
123456	300	800	160

Micro-batch
Output

Now the Store is updated

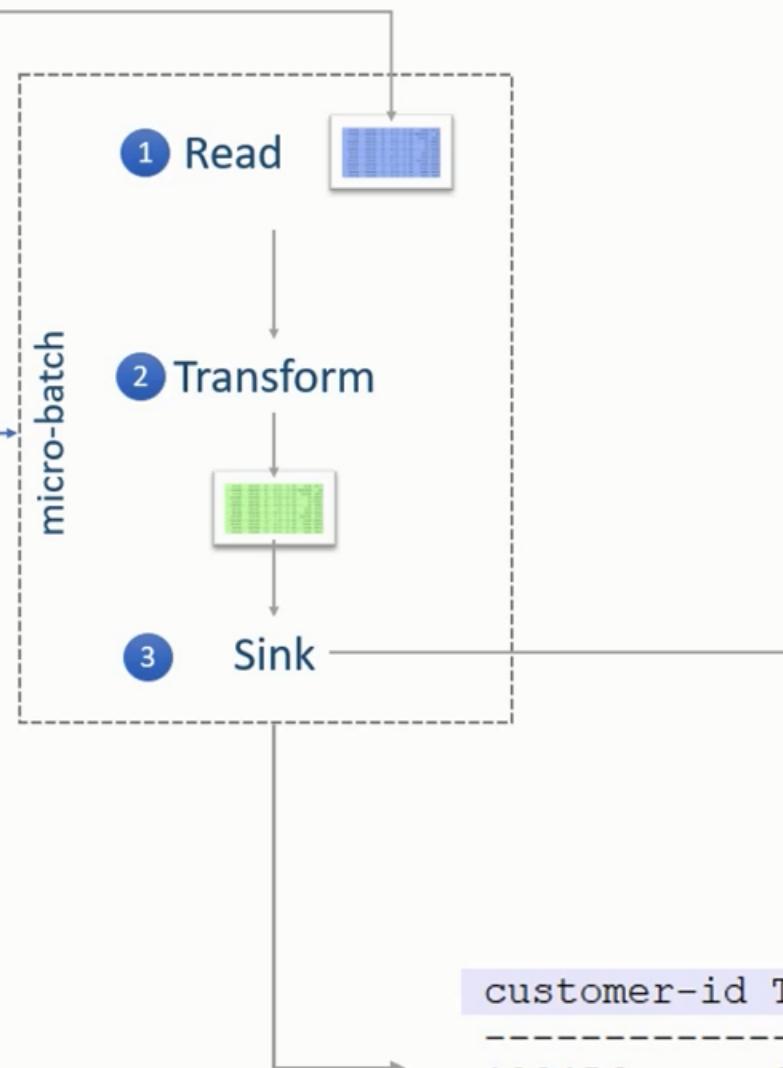
customer-id	Total Purchase	TotalRewards
123456	800	160

State Store



Streaming
Source

BG Thread



Third purchase

customer-id	Invoice Amount	Total Purchase	TotalRewards
123456	800	1600	320

Micro-batch
Output

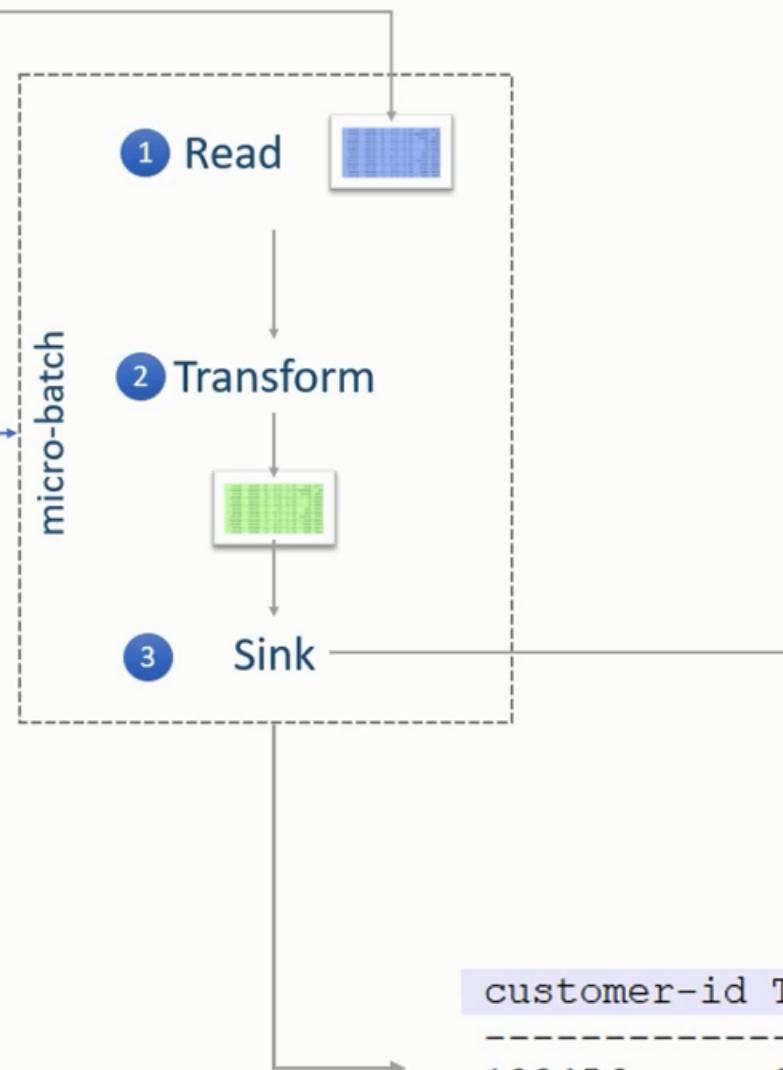
customer-id	Total Purchase	TotalRewards
123456	800	160

State Store



Streaming
Source

BG Thread



Third purchase

customer-id	Invoice Amount	Total Purchase	TotalRewards
123456	800	1600	320

Micro-batch
Output

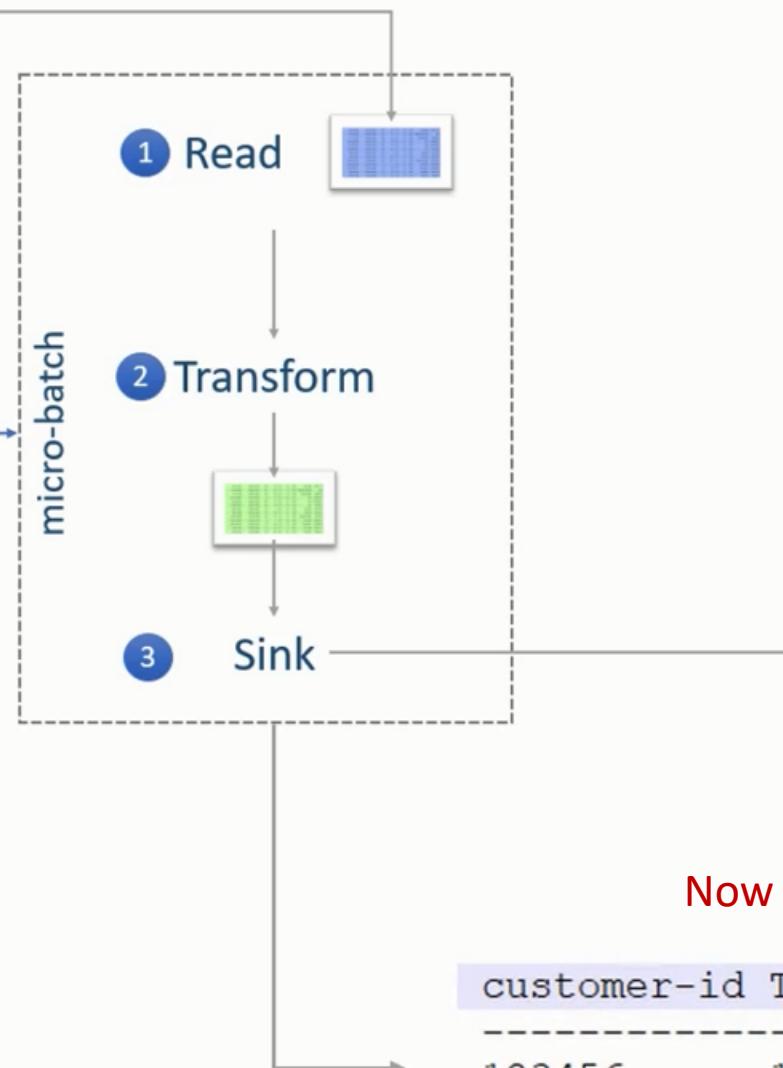
customer-id	Total Purchase	TotalRewards
123456	800	160

State Store



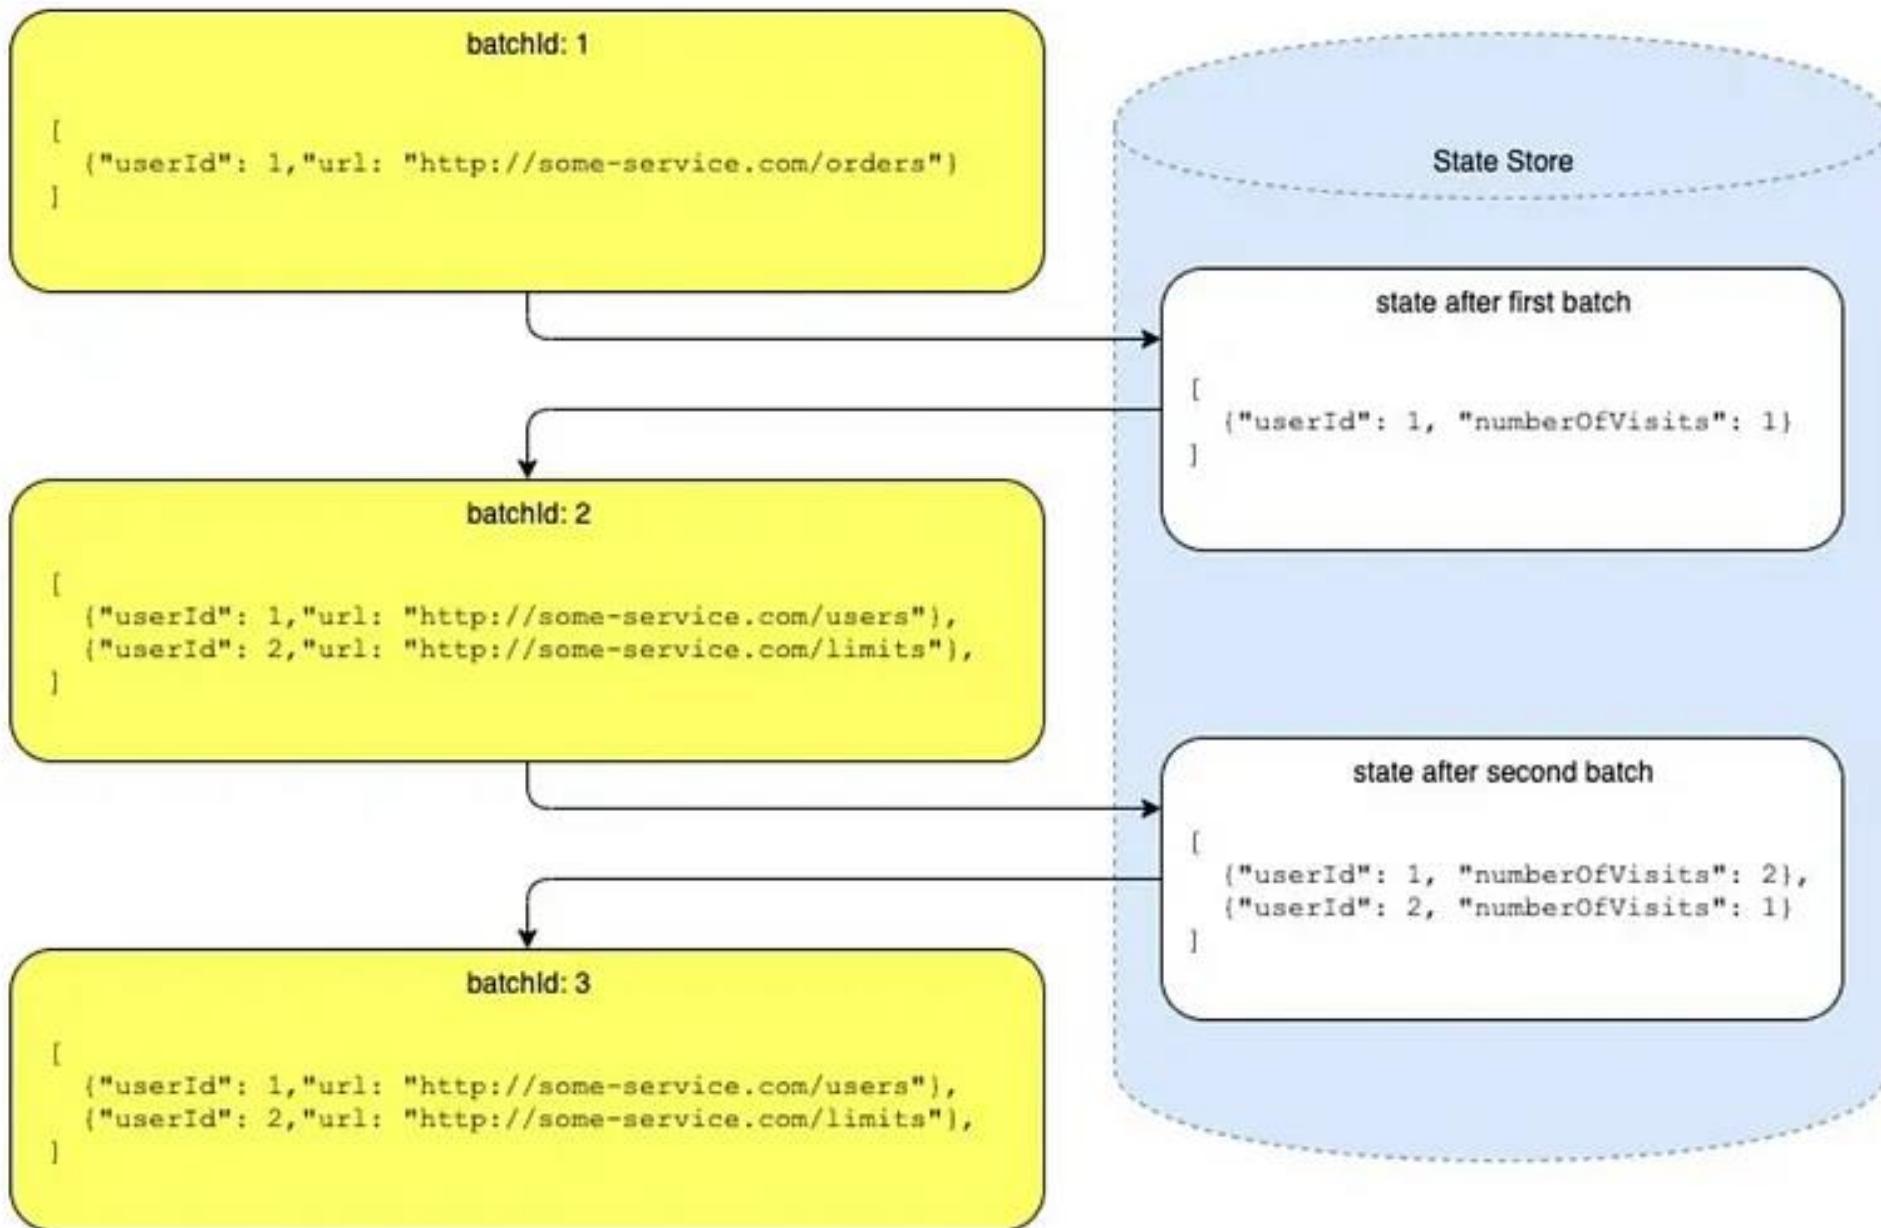
Streaming
Source

BG Thread



customer-id	Invoice Amount	Total Purchase	TotalRewards
123456	800	1600	320

Micro-batch
Output



State store and Spark Structured Streaming process schema



- 1. Stateless transformation**
select(), filter(), map(),
flatMap(), explode()
- 2. Stateful transformations**
Grouping, Aggregations
Windowing, and Joins



First we should **reflect on** these two categories.

1. **Stateless transformation**
Complete output mode not supported
2. **Stateful transformations**
Excessive state causes out of memory



Only supports "append" and
"update" mode!



1. Stateless transformation

Complete output mode not supported

2. Stateful transformations

Excessive state causes out of memory



1. **Stateless transformation**
Complete output mode not supported
2. **Stateful transformations**
Excessive state causes out of memory





When to use of **Managed** and
when use **Unmanaged**?

Stateful transformations

- Managed Stateful operations
- Unmanaged Stateful operations

Managed operation will be handle by Spark.

Unmanaged will be managed by user.

In the sequel, we just focus on Managed stateful operations.



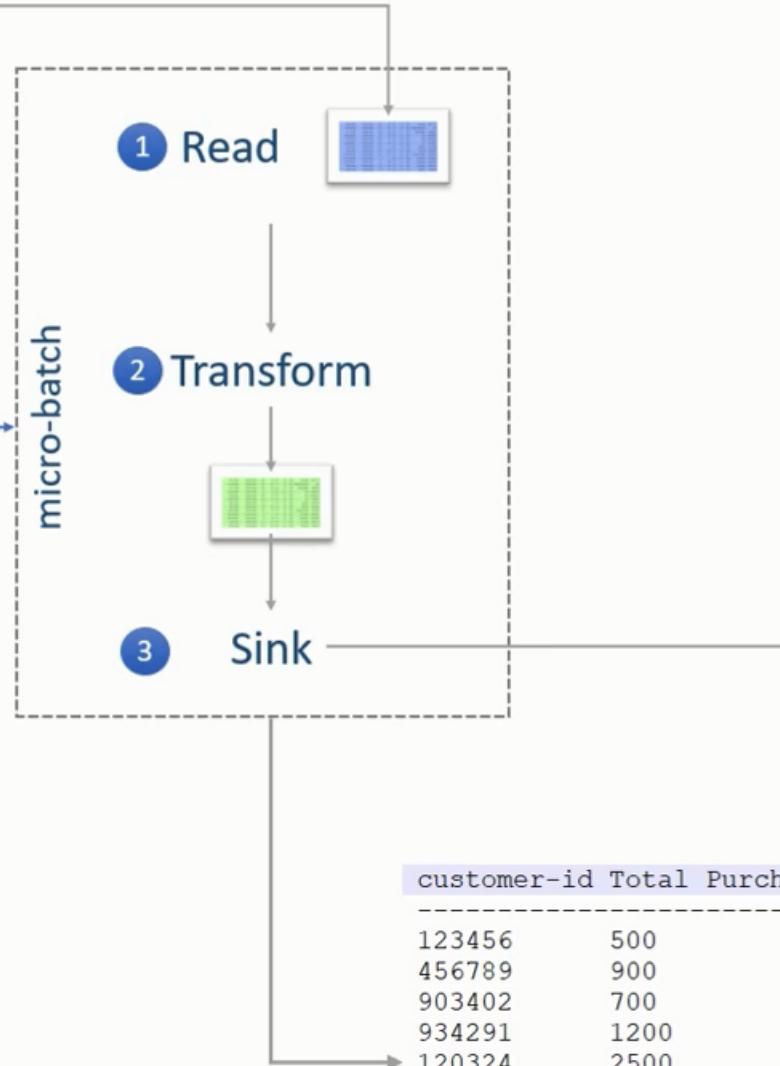
Aggregations

- Continuous aggregation
- Time-bound aggregations



Streaming
Source

BG Thread



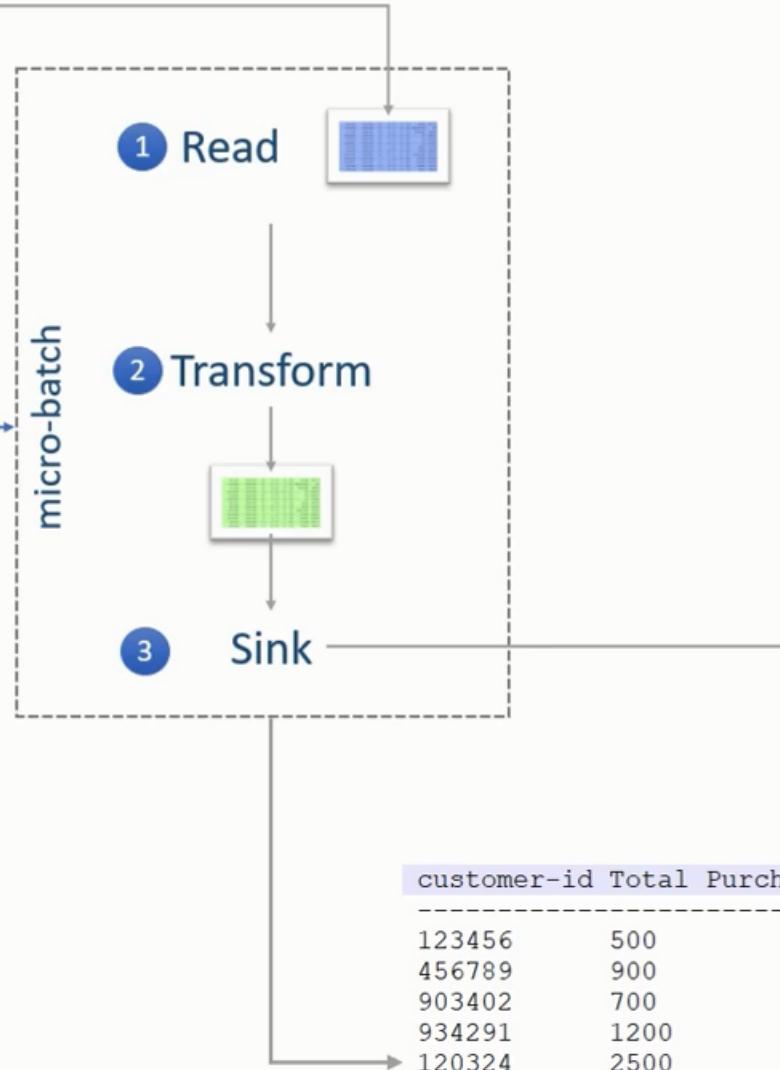
customer-id	Invoice Amount	Total Purchase	TotalRewards
123456	800	1600	320

Micro-batch
Output



Streaming
Source

BG Thread



State Store

customer-id	Total Purchase	TotalRewards	State Expiry Time
123456	500	100	18-Jul-2024
456789	900	180	
903402	700	140	
934291	1200	240	
120324	2500	500	
231095	100	20	
.....	

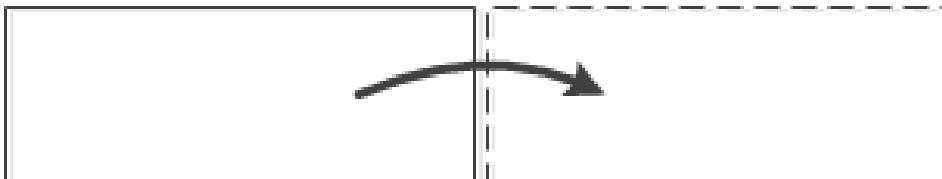
Windowing



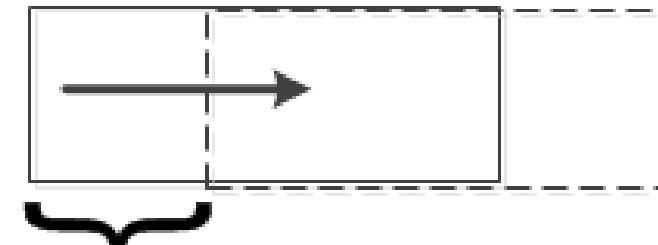
Windowing Aggregates

1. Tumbling Time Window
2. Sliding Time Window
3. Session Window

Tumbling Window
(Slide = WindowSize)



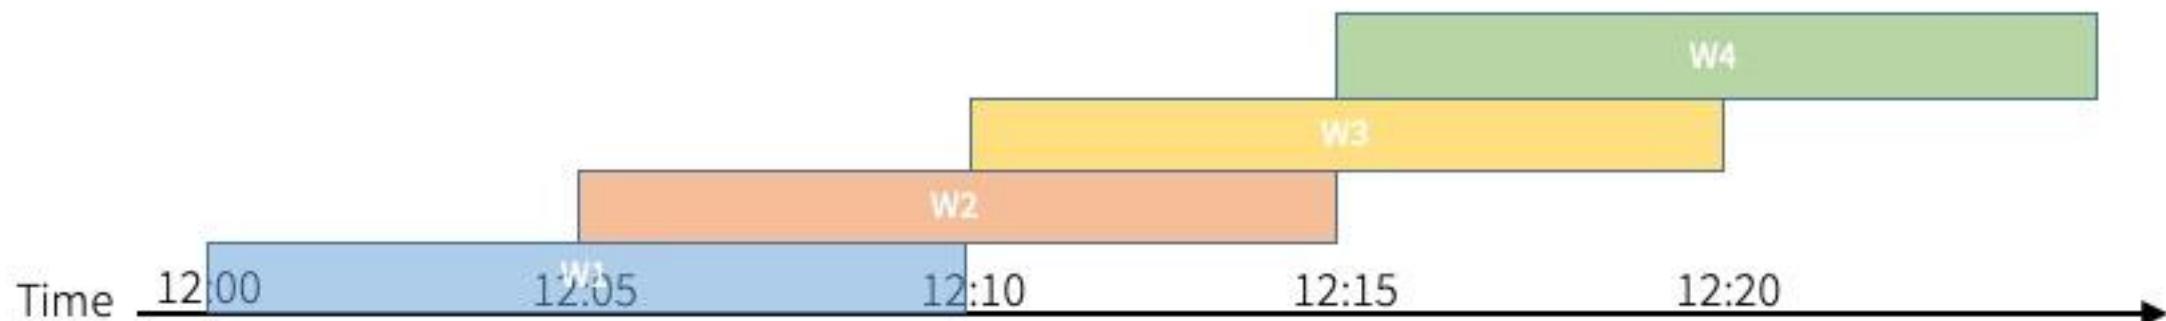
Sliding Window
(Slide < WindowSize)



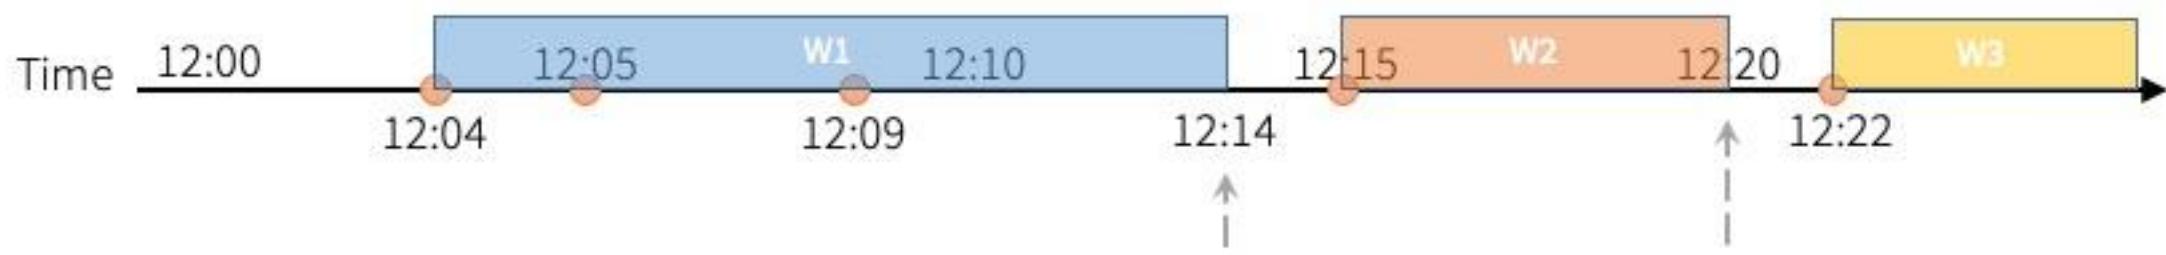
Tumbling Windows (5 mins)

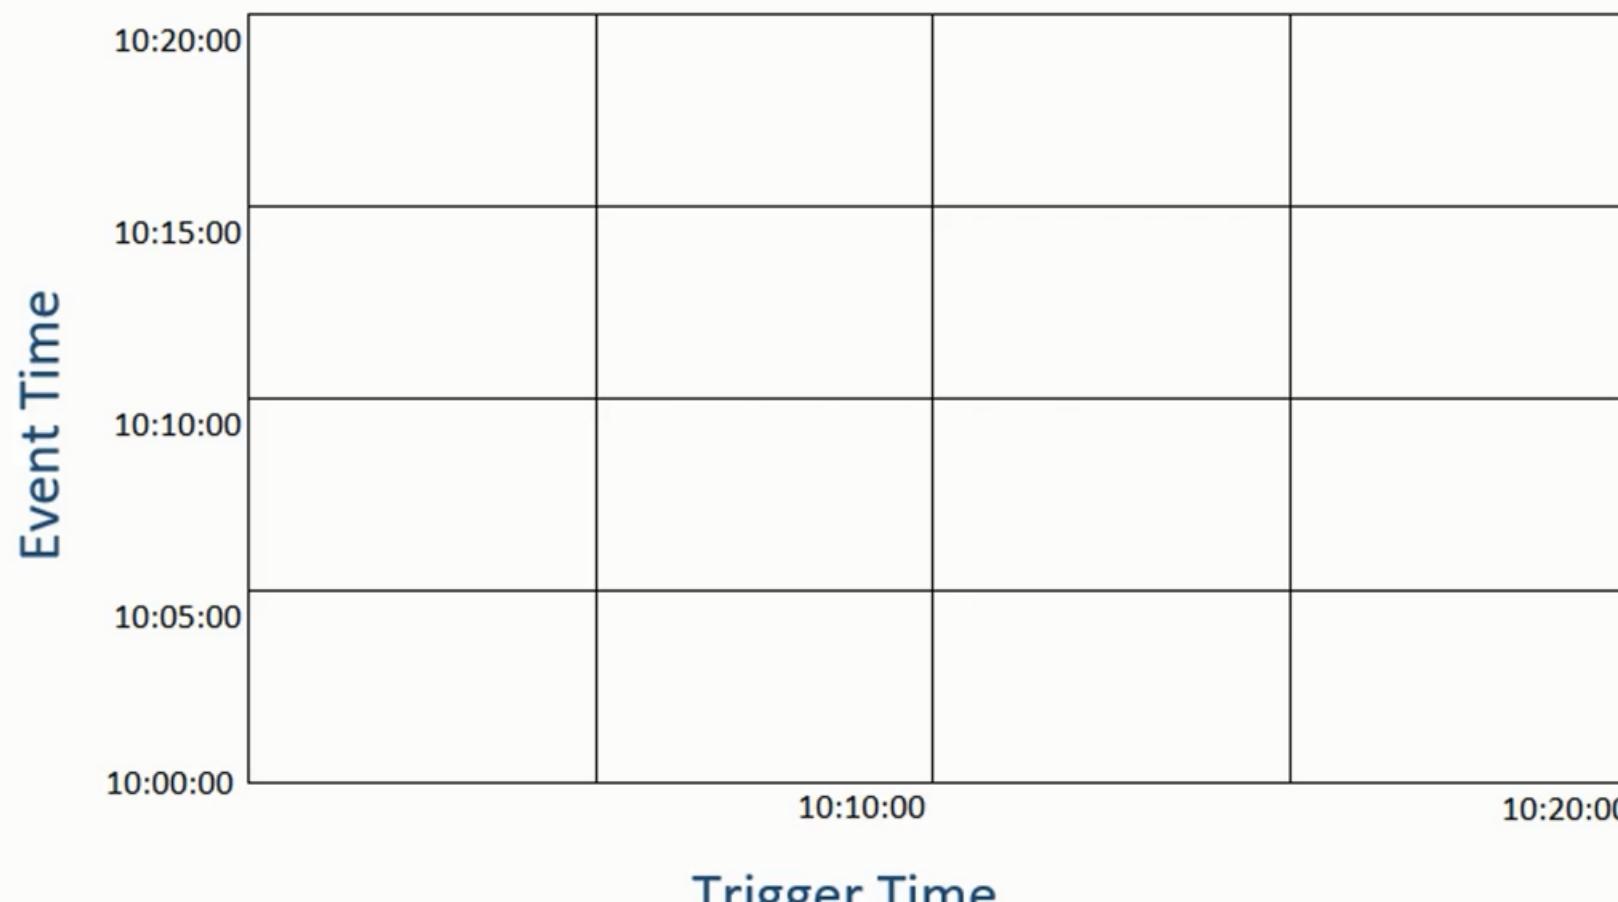


Sliding Windows (10 mins, slide 5 mins)

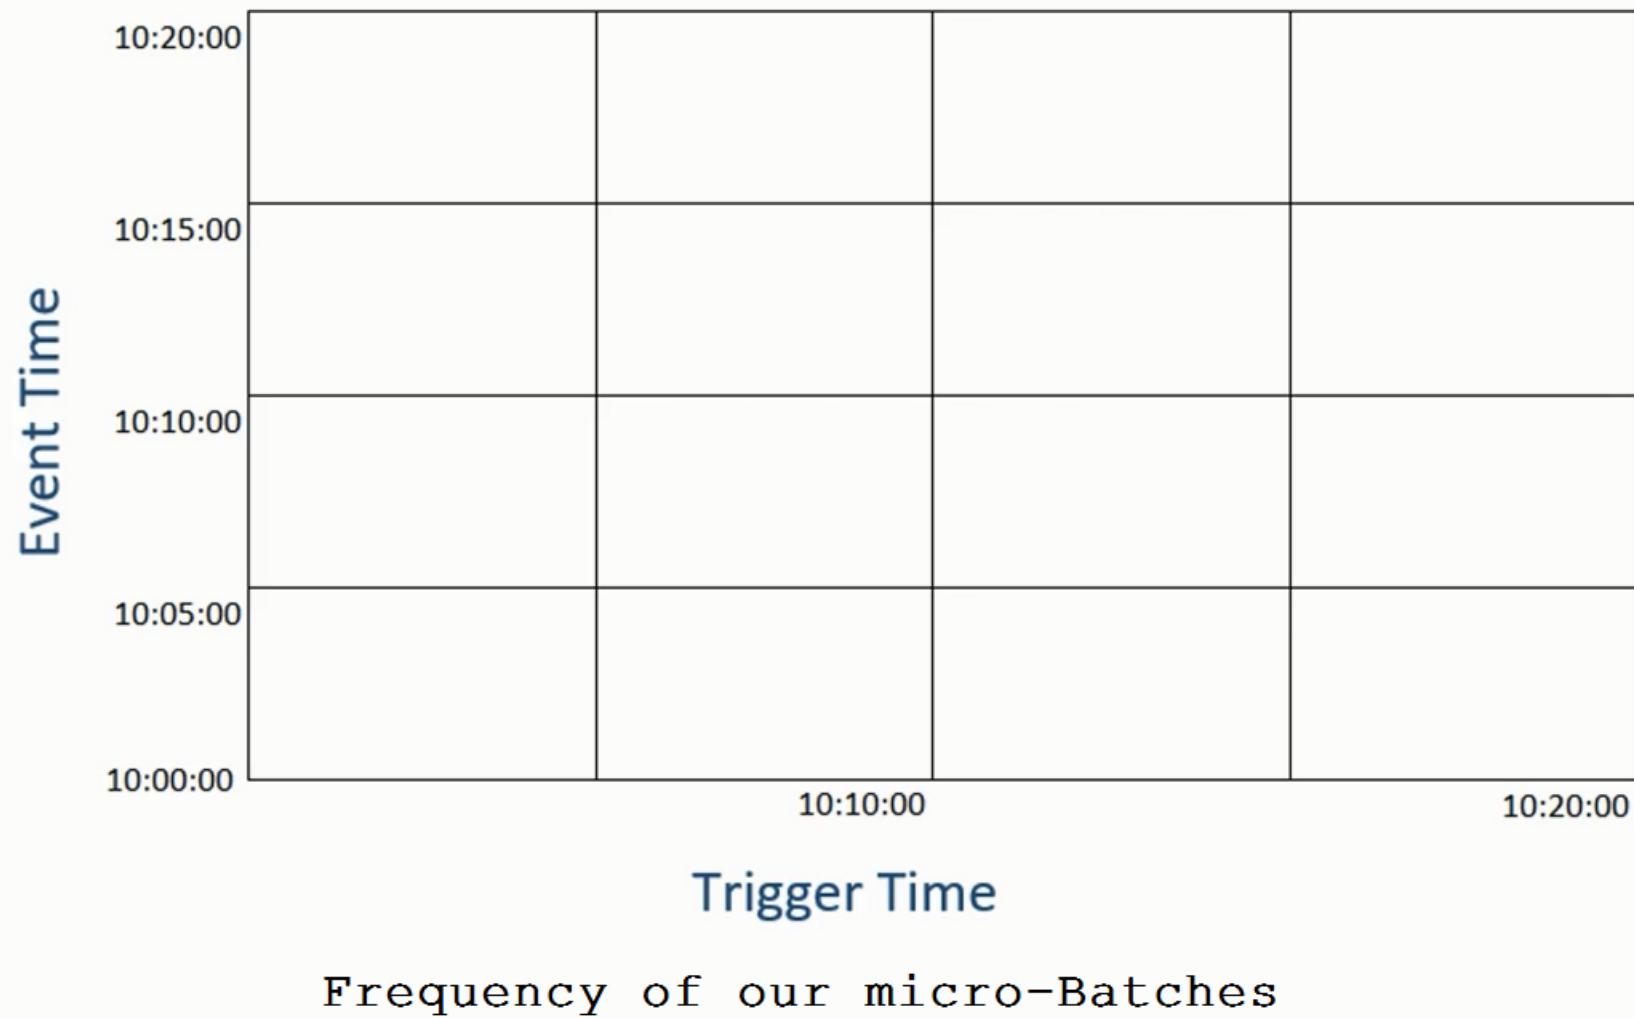


Session Windows (gap duration 5 mins)

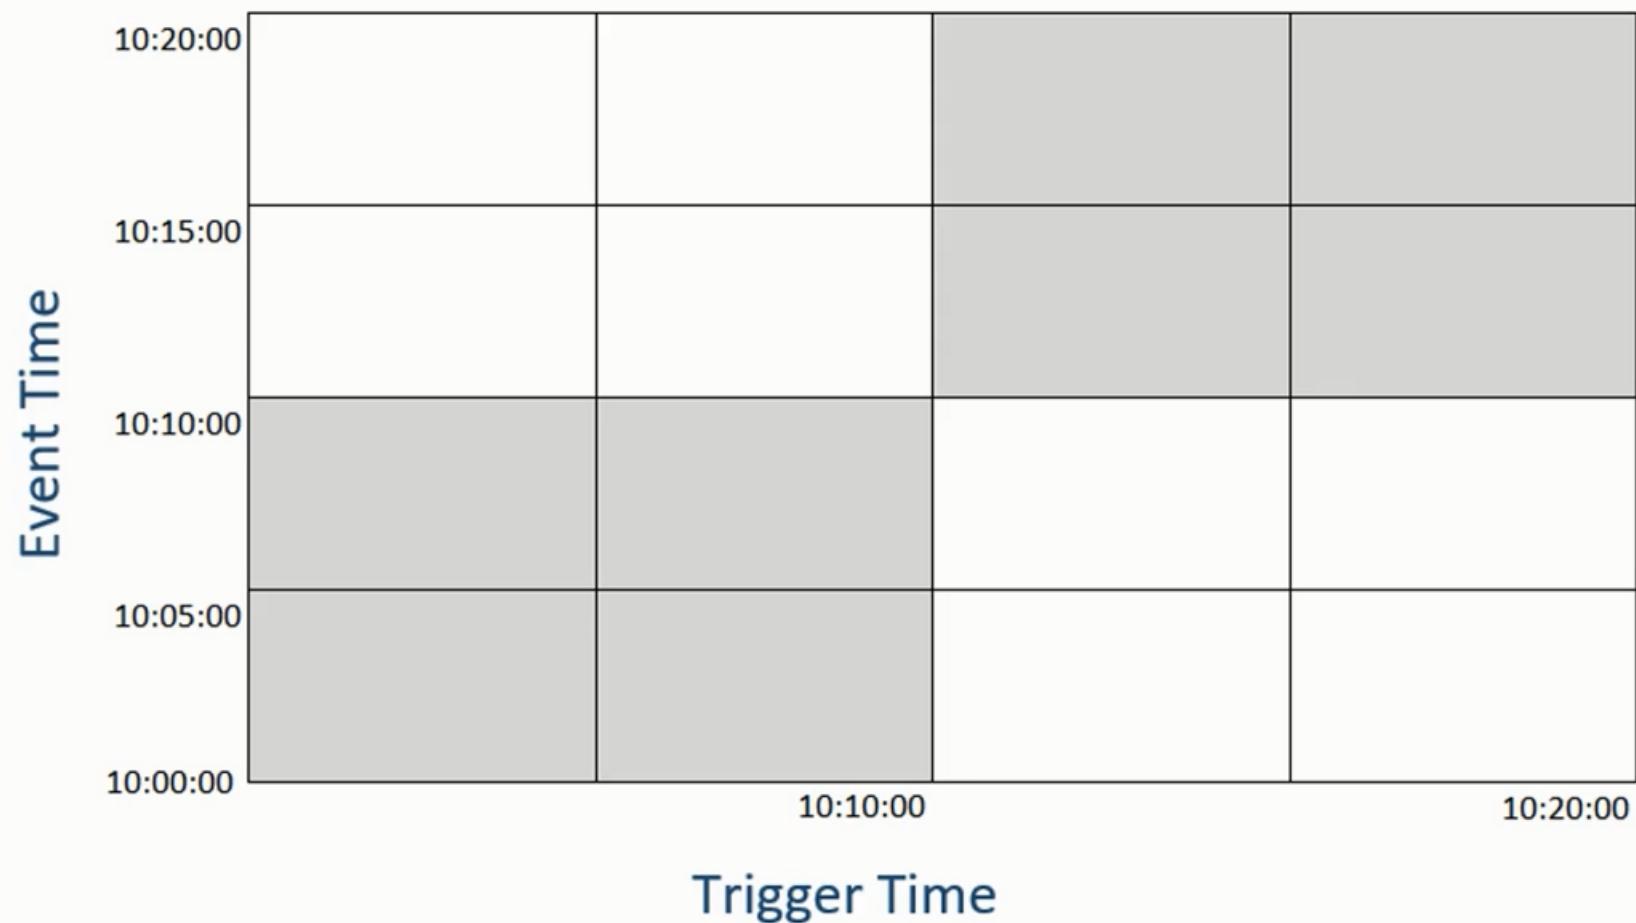




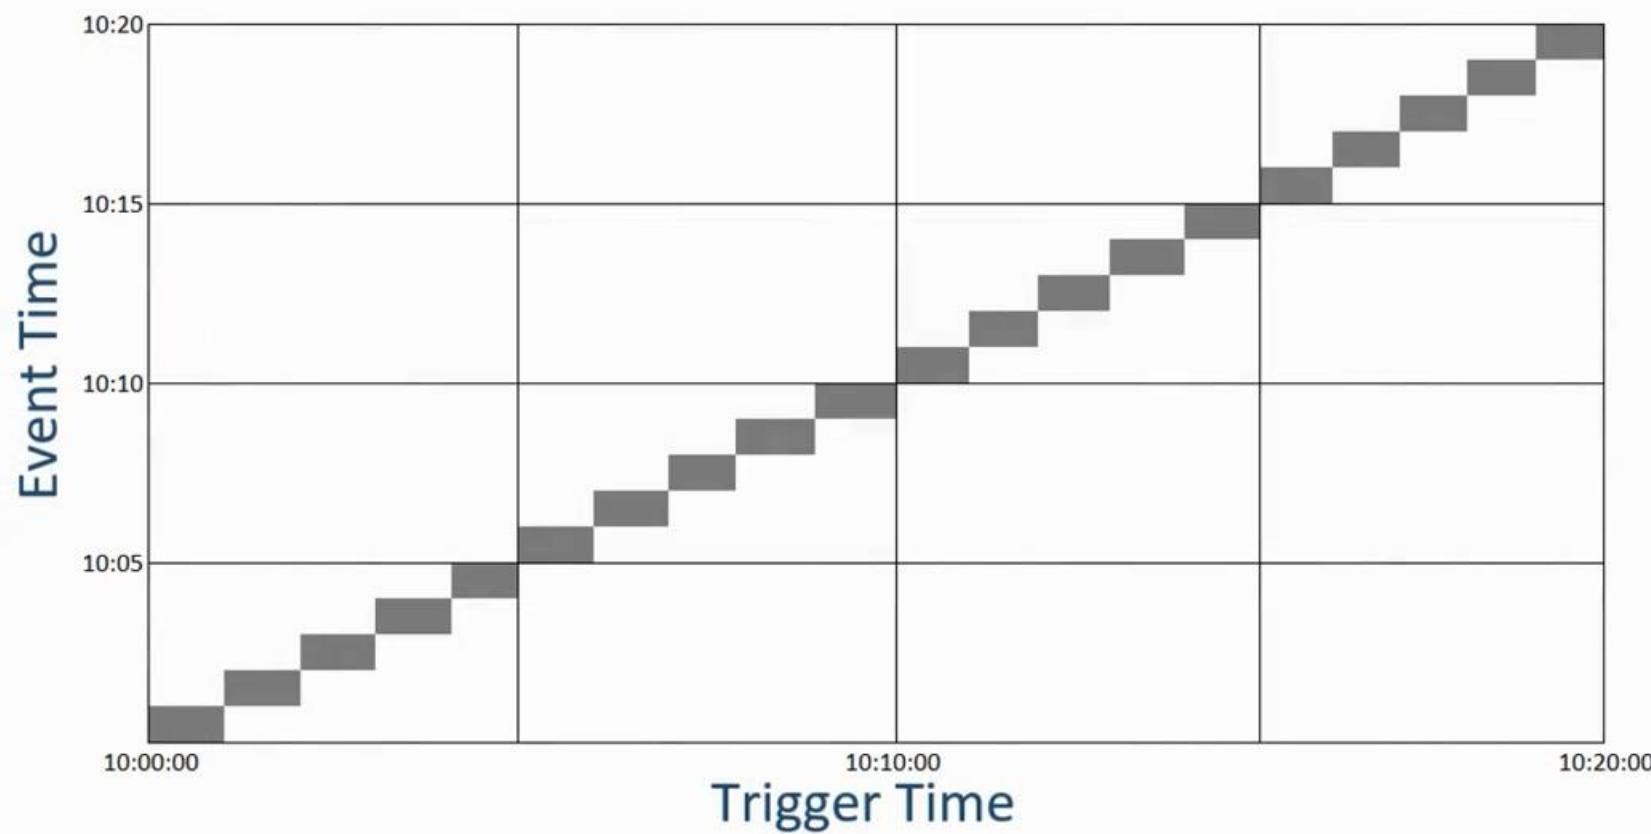
When our record
is generated

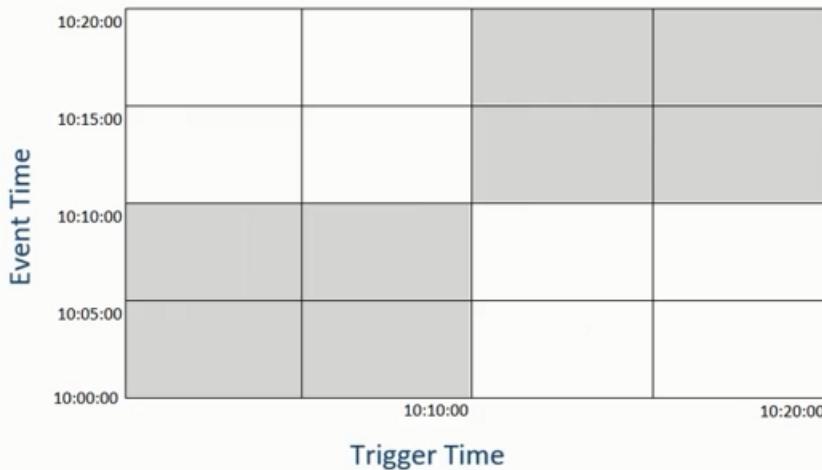
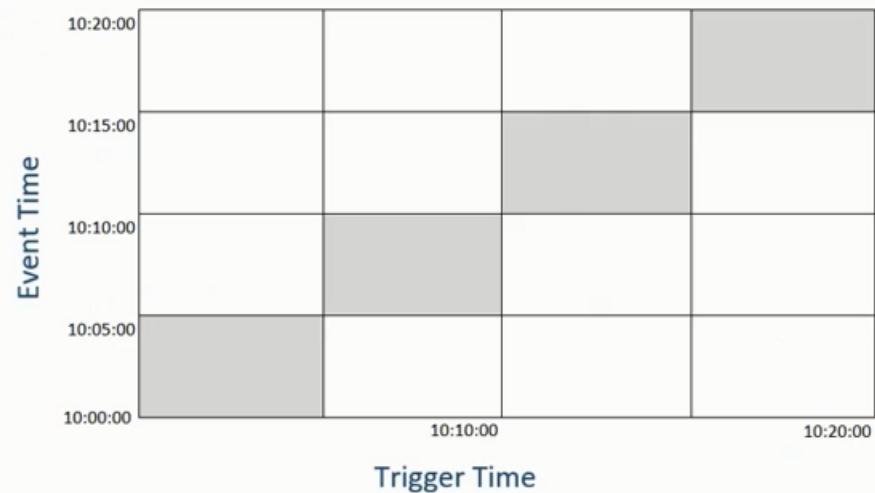
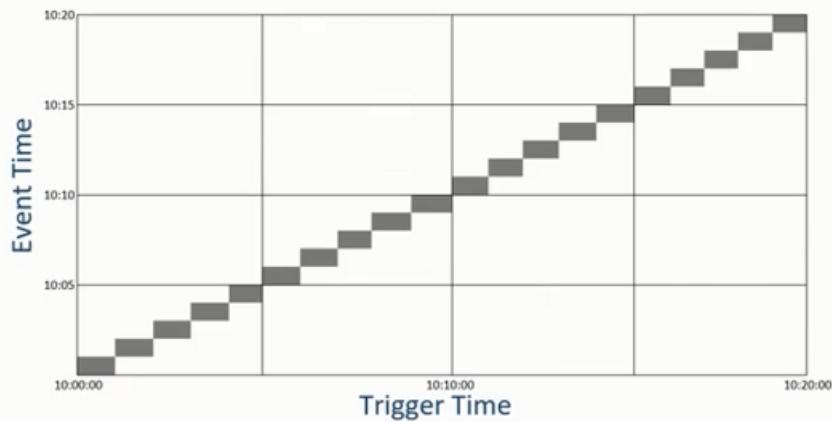






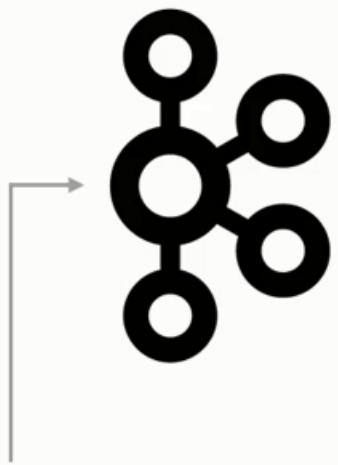






The Aggregation window has nothing to do with trigger time!

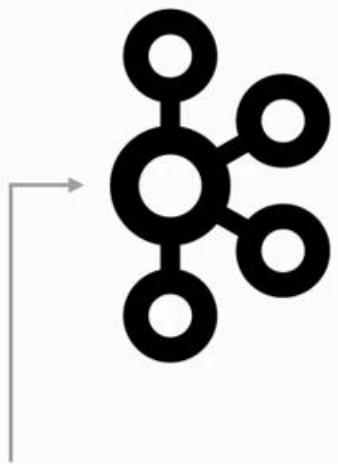
The Trigger is the time when we **start processing** our micro-batch



Trade Events

```
{"TransactionTime": "18:07:2024 10:01:12",
 "TransactionType": "BUY",
 "TransactionAmount": 500,
 "BrokerCode": "ABX"}
```

```
{"TransactionTime": "18:07:2024 10:05:30",
 "TransactionType": "SELL",
 "TransactionAmount": 300,
 "BrokerCode": "ABX"}
```



Trade Events

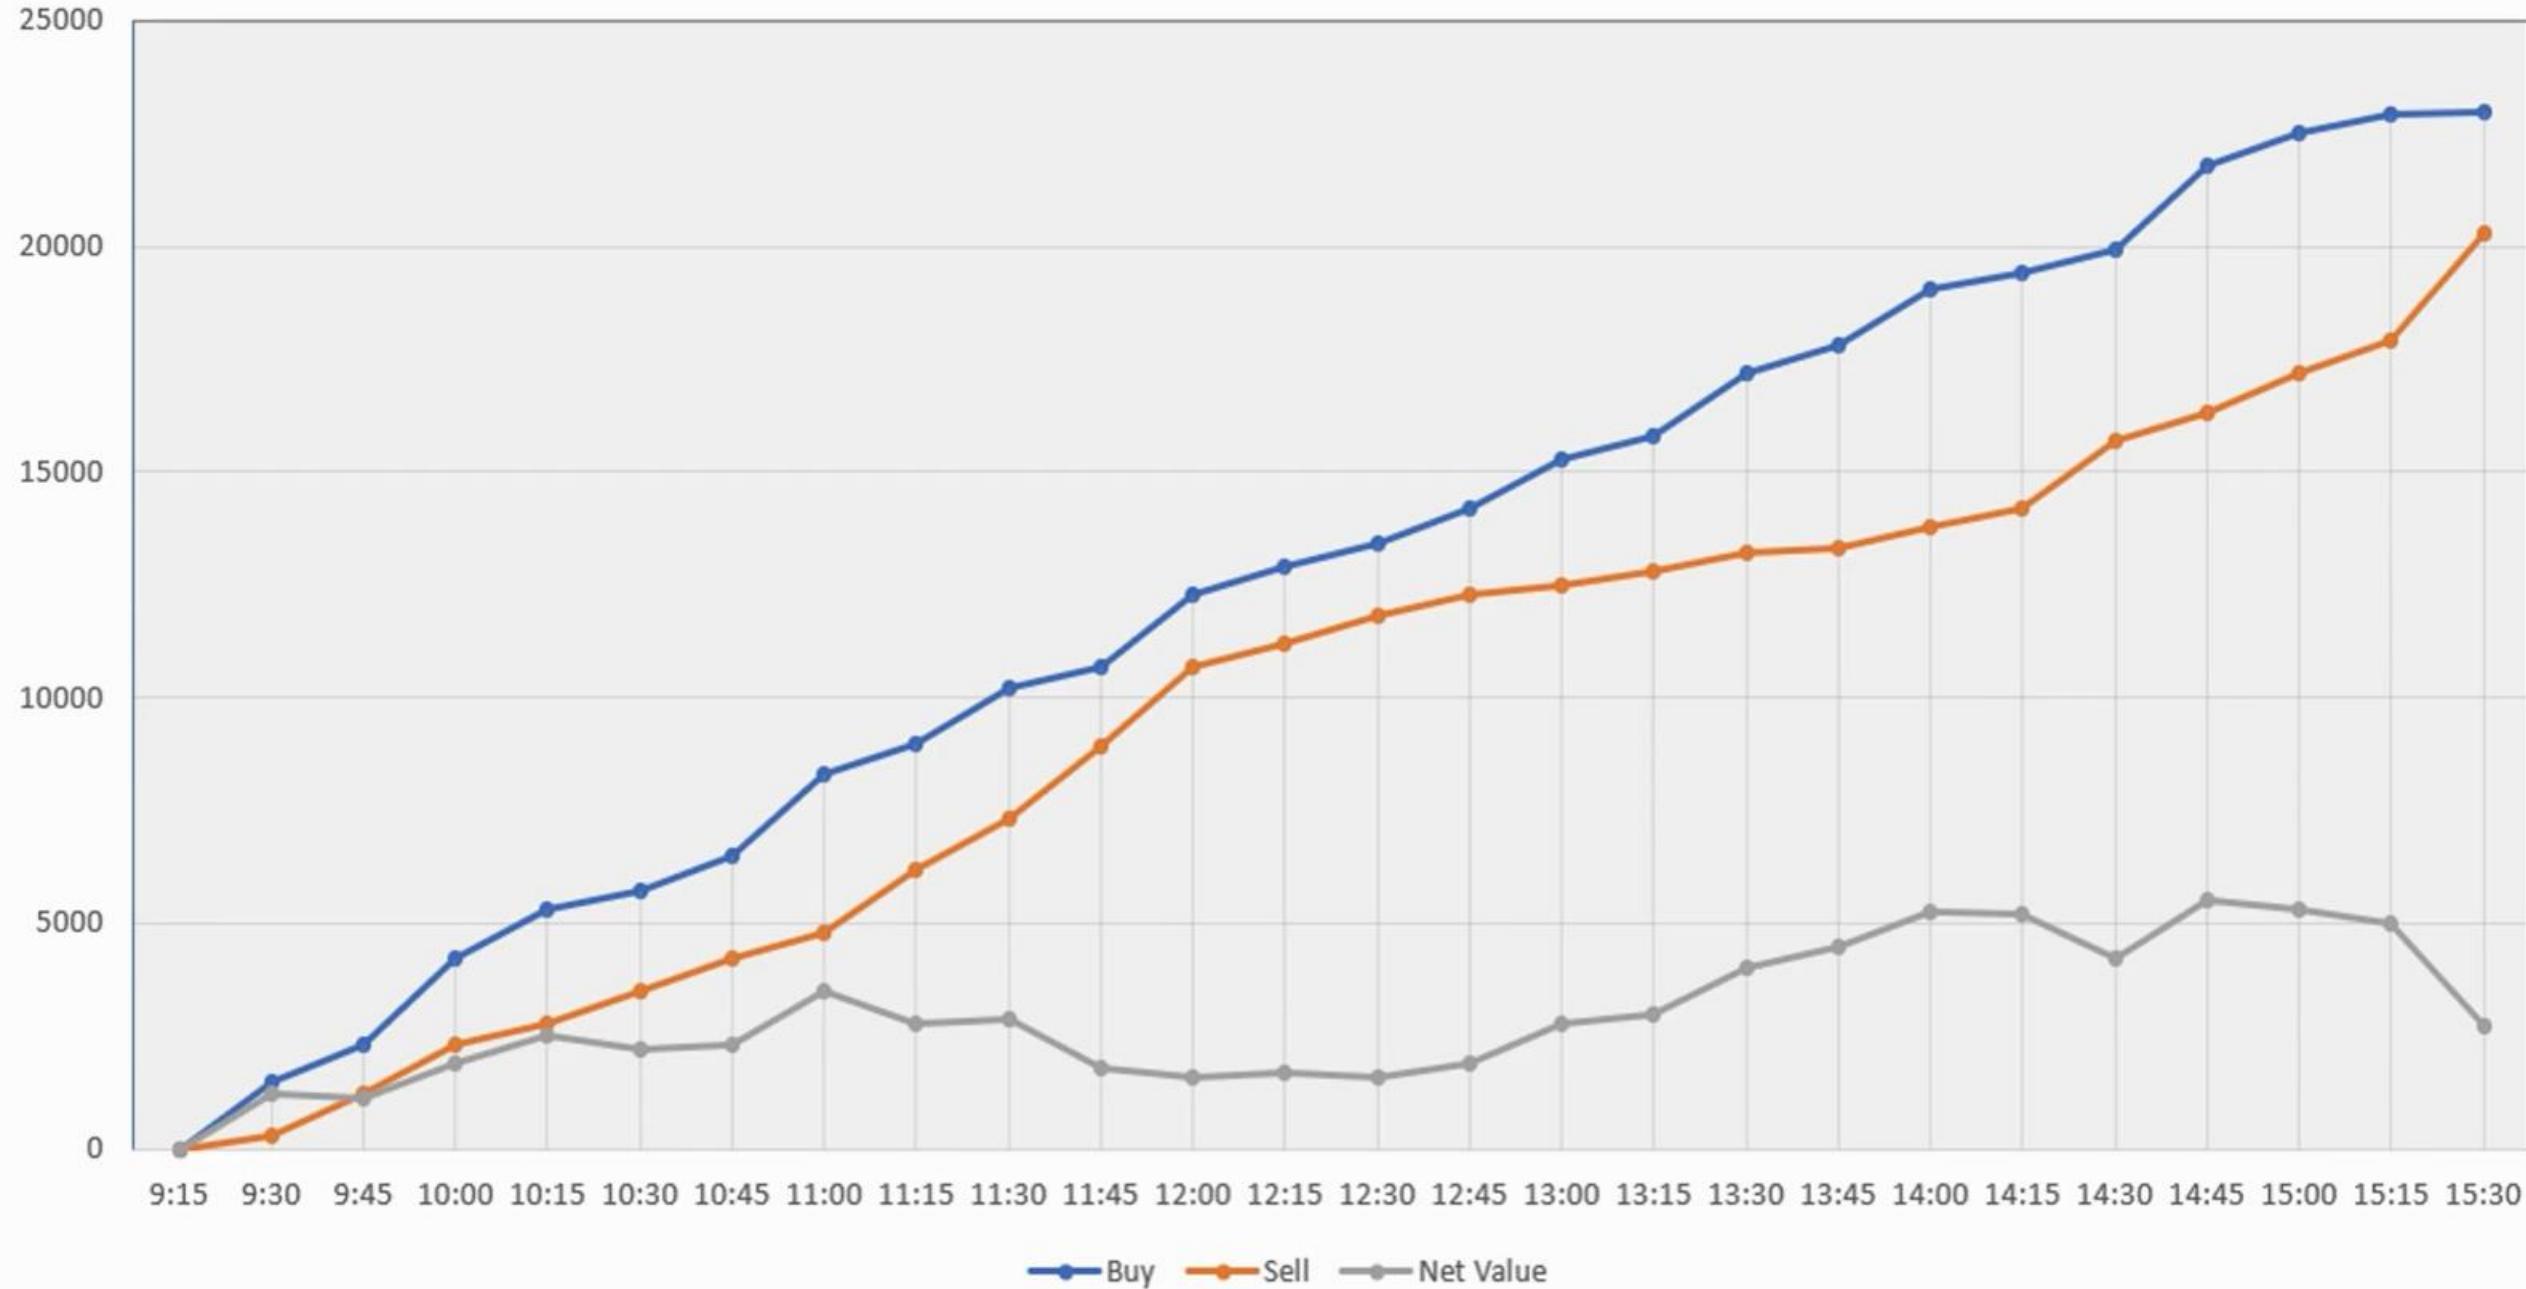
```
{"TransactionTime": "18:07:2024 10:01:12",
 "TransactionType": "BUY",
 "TransactionAmount": 500,
 "BrokerCode": "ABX"}
```

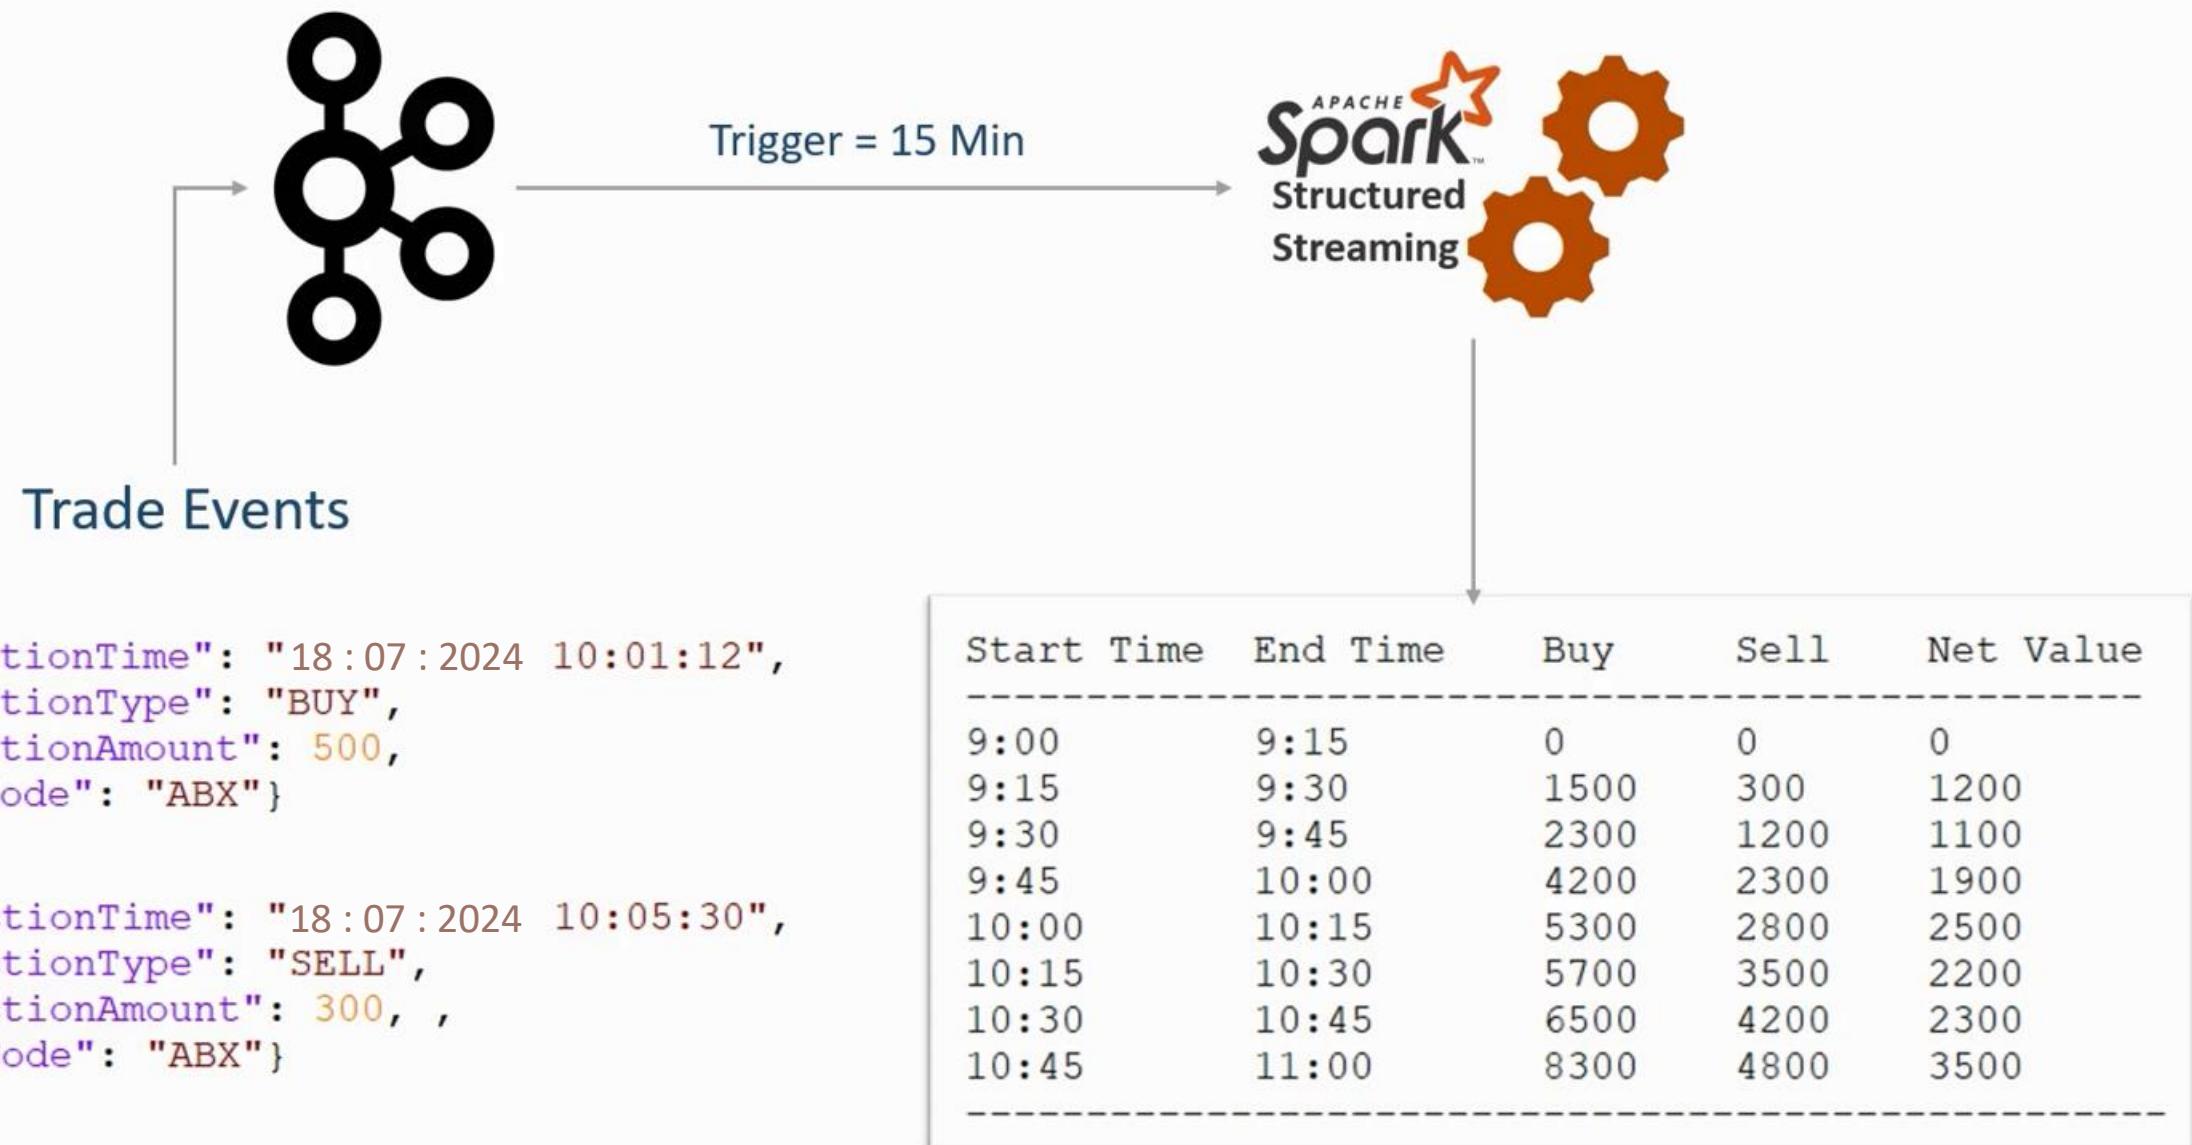
```
{"TransactionTime": "18:07:2024 10:05:30",
 "TransactionType": "SELL",
 "TransactionAmount": 300,
 "BrokerCode": "ABX"}
```



Start Time	End Time	Buy	Sell	Net Value
9:00	9:15	0	0	0
9:15	9:30	1500	300	1200
9:30	9:45	2300	1200	1100
9:45	10:00	4200	2300	1900
10:00	10:15	5300	2800	2500
10:15	10:30	5700	3500	2200
10:30	10:45	6500	4200	2300
10:45	11:00	8300	4800	3500

Day Trade

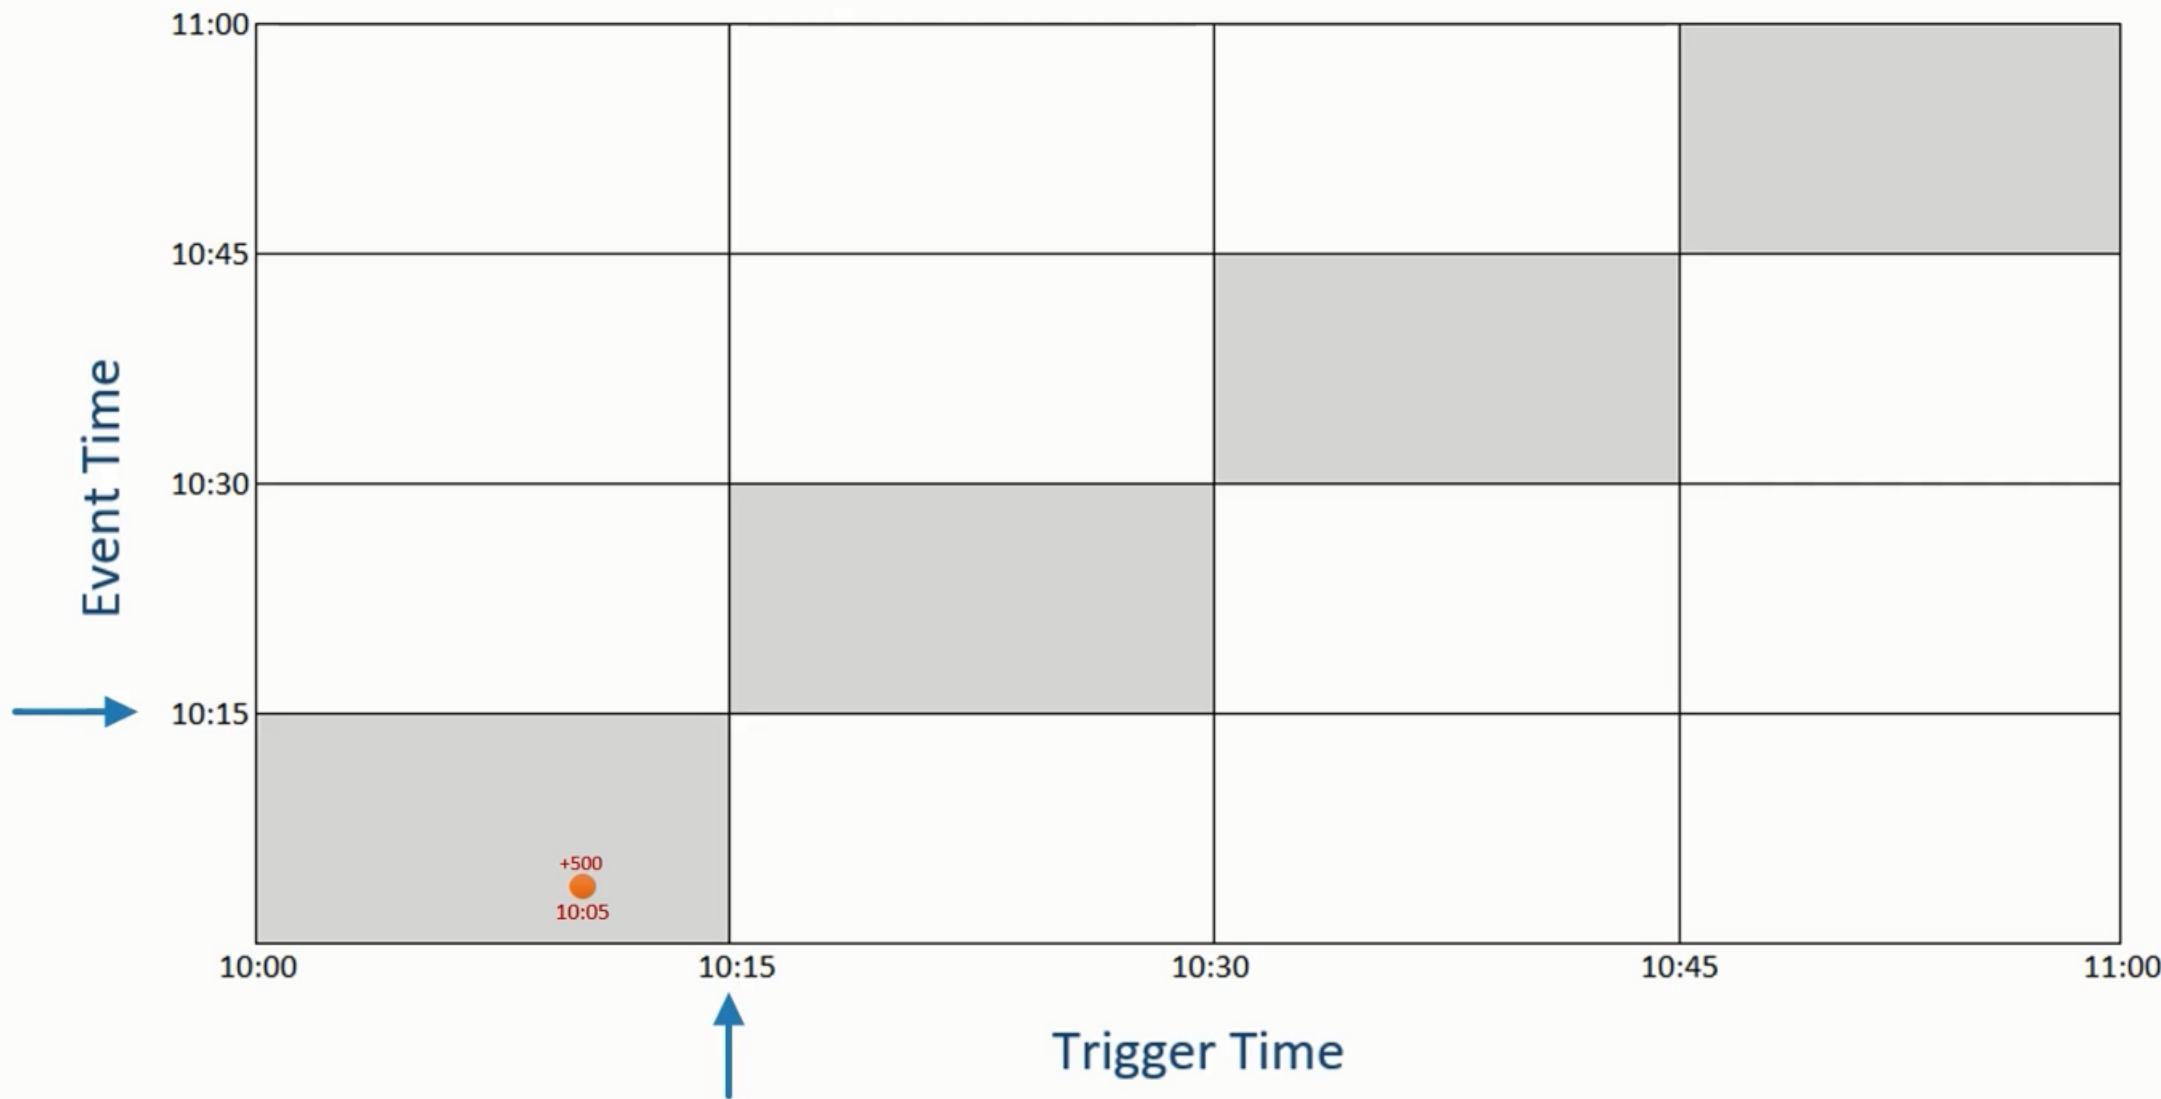


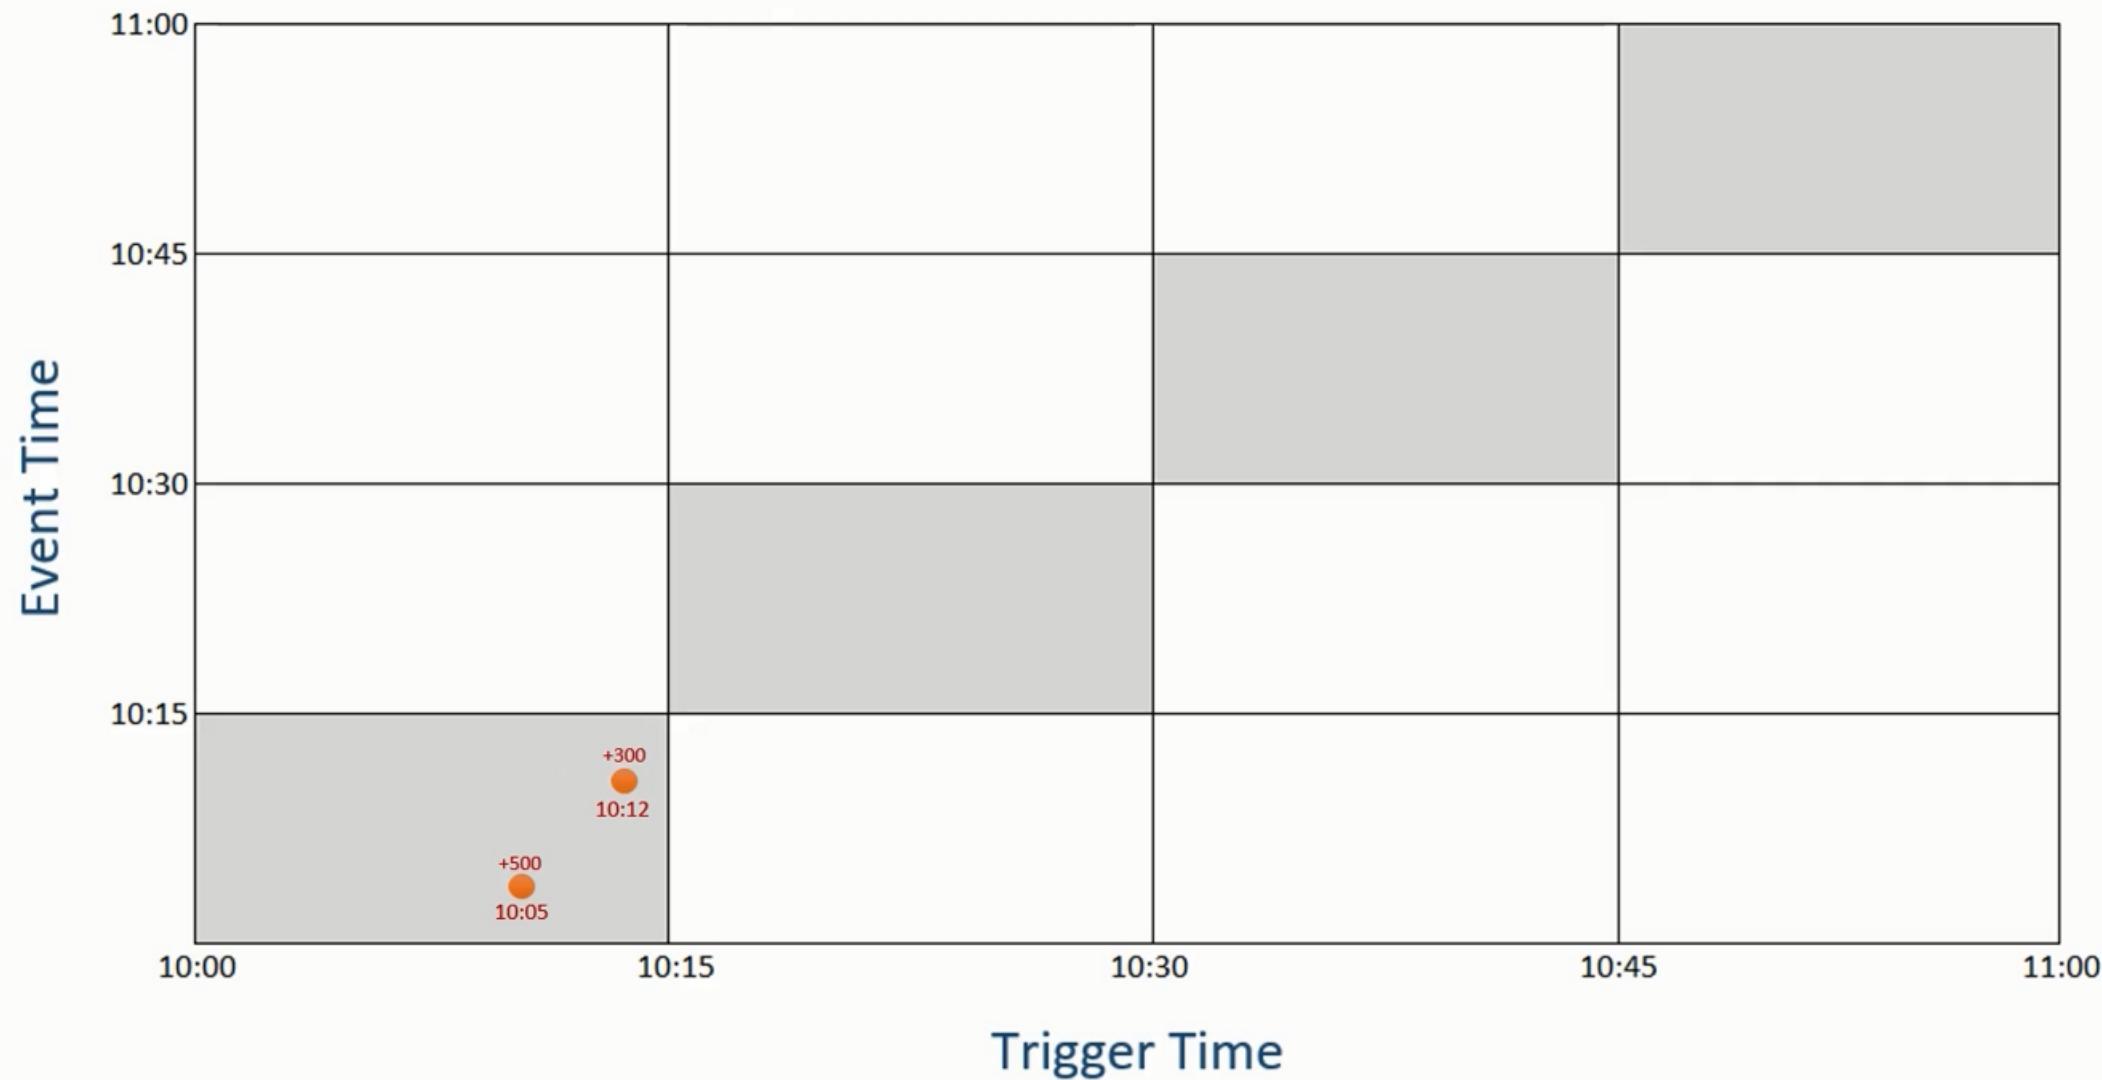




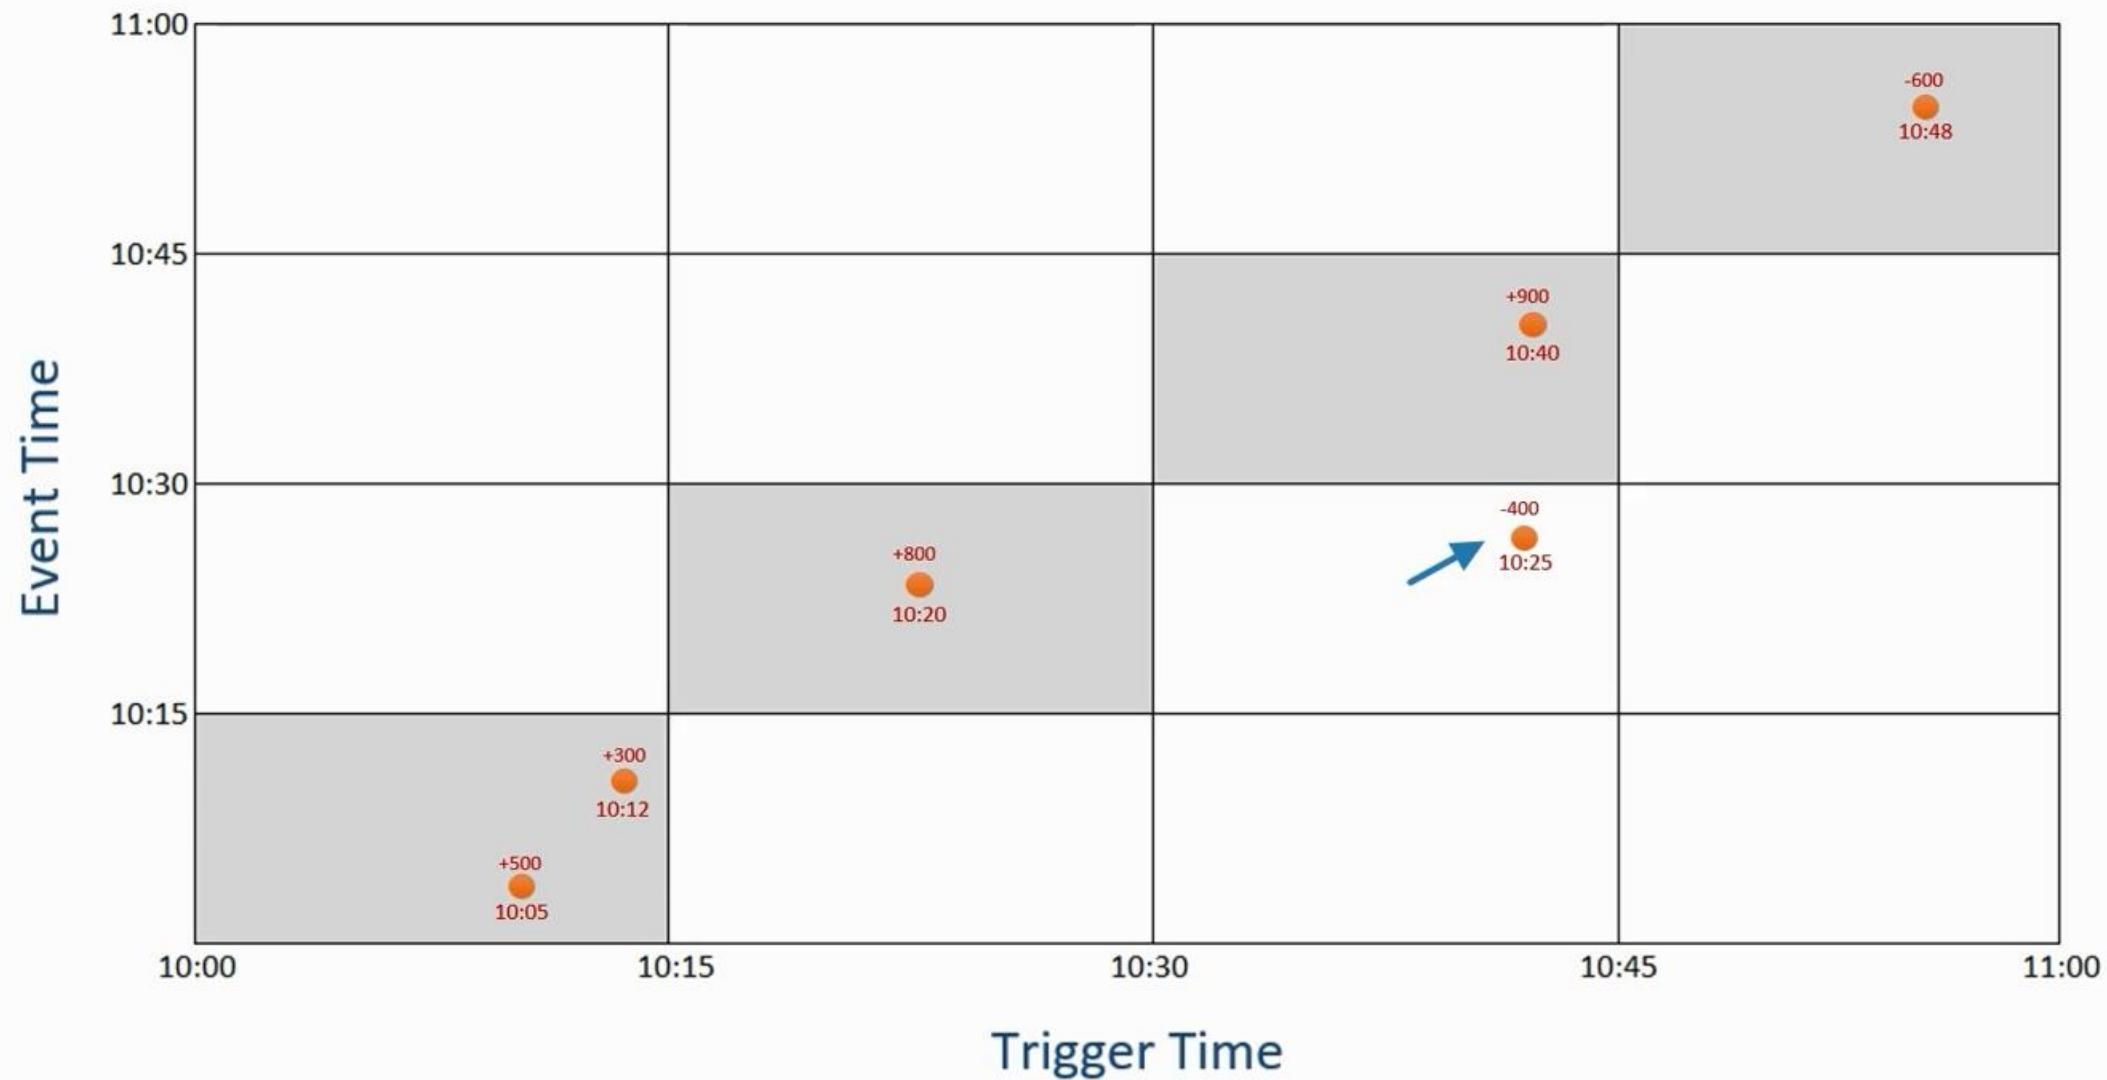
```
{"TransactionTime": "18:07:2024 10:01:12",
 "TransactionType": "BUY",
 "TransactionAmount": 500,
 "BrokerCode": "ABX"}  
  
{"TransactionTime": "18:07:2024 10:05:30",
 "TransactionType": "SELL",
 "TransactionAmount": 300,
 "BrokerCode": "ABX"}
```

Start Time	End Time	Buy	Sell	Net Value
9:00	9:15	0	0	0
9:15	9:30	1500	300	1200
9:30	9:45	2300	1200	1100
9:45	10:00	4200	2300	1900
10:00	10:15	5300	2800	2500
10:15	10:30	5700	3500	2200
10:30	10:45	6500	4200	2300
10:45	11:00	8300	4800	3500









Now let's go for coding!